

Keras -- MLPs on MNIST

In [18]:

```
# if you keras is not using tensorflow as backend set "KERAS_
BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import matplotlib.pyplot as plt
```

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25
# 362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [20]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [21]:

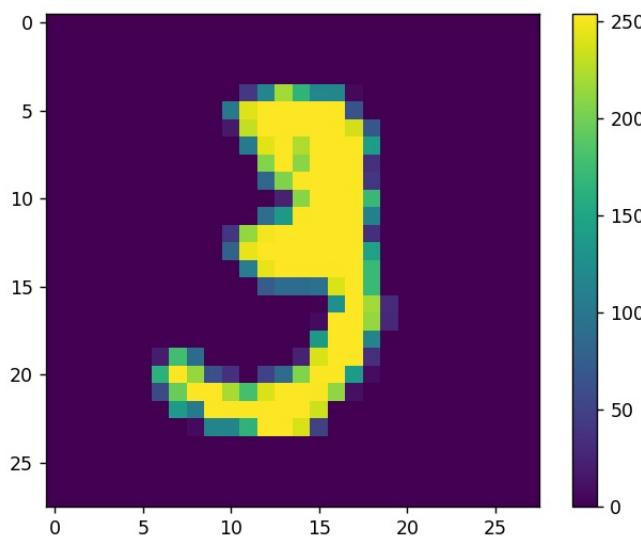
```
print("Number of training examples :", X_train.shape[0], "and  
each image is of shape (%d, %d)"%(X_train.shape[1], X_train.  
shape[2]))  
print("Number of testing examples :", X_test.shape[0], "and e  
ach image is of shape (%d, %d)"%(X_test.shape[1], X_test.shap  
e[2]))
```

Number of training examples : 60000 and each i
mage is of shape (28, 28)

Number of testing examples : 10000 and each im
age is of shape (28, 28)

In [22]:

```
#viewing a sample data point from training  
plt.figure()  
plt.imshow(X_train[10])  
plt.colorbar()  
plt.grid(False)  
plt.show()
```



In [23]:

```
#normalising the data# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255

X_train = X_train/255
X_test = X_test/255
```

In [24]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [25]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of testing examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of testing examples : 10000 and each image is of shape (784)

In [26]:

```
x_train[0]
```

Out[26]:

6862745, 0.60392157,
0.66666667, 0.99215686, 0.99215686, 0.9
9215686, 0.99215686,
0.99215686, 0.88235294, 0.6745098 , 0.9
9215686, 0.94901961,
0.76470588, 0.25098039, 0. . , 0.
. , 0. . , 0. . , 0. .
. , 0. . , 0. . , 0. .
0. . , 0. . , 0. . , 0.1
9215686, 0.93333333,
0.99215686, 0.99215686, 0.99215686, 0.9
9215686, 0.99215686,
0.99215686, 0.99215686, 0.99215686, 0.9
8431373, 0.36470588,
0.32156863, 0.32156863, 0.21960784, 0.1
5294118, 0. . ,
0. . , 0. . , 0. . , 0. .
. , 0. . , 0. . , 0. .
0. . , 0. . , 0. . , 0. .
. , 0. . , 0. . , 0. .
0. . , 0.07058824, 0.85882353, 0.9
9215686, 0.99215686,
0.99215686, 0.99215686, 0.99215686, 0.7
7647059, 0.71372549,
0.96862745, 0.94509804, 0. . , 0.
. , 0. . ,
0. . , 0. . , 0. . , 0.
. , 0. . ,
0. . , 0. . , 0. . , 0.
. , 0. . ,
0. . , 0. . , 0. . , 0.
. , 0. . ,
0.31372549, 0.61176471, 0.41960784, 0.9
9215686, 0.99215686,
0.80392157, 0.04313725, 0. . , 0.1
6862745, 0.60392157,

0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0.05490196,
0.00392157, 0.60392157, 0.99215686, 0.3
5294118, 0. ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0.54509804,
0.99215686, 0.74509804, 0.00784314, 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0.04313725, 0.7
4509804, 0.99215686,
0.2745098 , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.
, 0. , ,
0. , 0. , 0. , 0.

, 0. , , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0.1372549 , 0.94509804, 0.8
8235294, 0.62745098,
0.42352941, 0.00392157, 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0.31764706, 0.94117647, 0.99215686, 0.9
9215686, 0.46666667,
0.09803922, 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0.17647059,
0.72941176, 0.99215686, 0.99215686, 0.5
8823529, 0.10588235,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,
0. , 0. , 0. , 0.
, 0. , , ,

0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.0
627451 , 0.36470588,
0.98823529, 0.99215686, 0.73333333, 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.9
7647059 , 0.99215686,
0.97647059, 0.25098039, 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.1
8039216 , 0.50980392,
0.71764706, 0.99215686, 0.99215686, 0.8
1176471 , 0.00784314,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0.15294118,
0.58039216, 0.89803922, 0.99215686, 0.9

9215686, 0.99215686,
0.98039216, 0.71372549, 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0.09411765, 0.44705882, 0.86666667, 0.9
9215686, 0.99215686,
0.99215686, 0.99215686, 0.78823529, 0.3
0588235, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0.09019608, 0.25882353, 0.8
3529412, 0.99215686,
0.99215686, 0.99215686, 0.99215686, 0.7
7647059, 0.31764706,
0.00784314, 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0.07058824, 0.6
7058824, 0.85882353,
0.99215686, 0.99215686, 0.99215686, 0.9
9215686, 0.76470588,
0.31372549, 0.03529412, 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,


```
, 0.
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
, 0. ,
0. , 0. , 0. , 0.
])
```

In [27]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0,
# , 0, 1, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

In [28]:

```
# https://keras.io/getting-started/sequential-model-guide/  
  
# The Sequential model is a linear stack of layers.  
# you can create a Sequential model by passing a list of layer instances to the constructor:  
  
# model = Sequential([  
#     Dense(32, input_shape=(784,)),  
#     Activation('relu'),  
#     Dense(10),  
#     Activation('softmax'),  
# ])  
  
# You can also simply add layers via the .add() method:  
  
# model = Sequential()  
# model.add(Dense(32, input_dim=784))  
# model.add(Activation('relu'))  
  
###  
  
# https://keras.io/layers/core/  
  
# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',  
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,  
# activity_regularizer=None,  
# kernel_constraint=None, bias_constraint=None)  
  
# Dense implements the operation: output = activation(dot(inp
```

```
ut, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

#####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [29]:

```
# some model parameters
```

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

In [30]:

```
# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic
# shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of
# input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

WARNING:tensorflow:From C:\Users\Public\Anaconda3\lib\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.

In [31]:

```
# Before training a model, you need to configure the learning
process, which is done via the compile method
```

```
# It receives three arguments:  
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/  
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/  
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/  
  
# Note: when using the categorical_crossentropy loss, your targets should be in categorical format  
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except  
# for a 1 at the index corresponding to the class of the sample).  
  
# that is why we converted our labels into vectors  
  
model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])  
  
# Keras models are trained on Numpy arrays of input data and labels.  
# For training a model, you will typically use the fit function  
  
# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_split=0.0,  
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,  
# validation_steps=None)  
  
# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).  
  
# it returns A History object. Its History.history attribute
```

is a record of training loss values and # metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

```
# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size,
epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From C:\Users\Public\Anaconda3\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

```
60000/60000 [=====] -  
 1s 13us/step - loss: 1.2957 - acc: 0.6885 - val_loss: 0.8133 - val_acc: 0.8292
```

Epoch 2/20

```
60000/60000 [=====] -  
 1s 12us/step - loss: 0.7187 - acc: 0.8399 - val_loss: 0.6077 - val_acc: 0.8601
```

Epoch 3/20

```
60000/60000 [=====] -  
 1s 11us/step - loss: 0.5881 - acc: 0.8593 - val_loss: 0.5249 - val_acc: 0.8751
```

Epoch 4/20

```
60000/60000 [=====] -  
 1s 11us/step - loss: 0.5256 - acc: 0.8689 - val_loss: 0.4790 - val_acc: 0.8828
```

Epoch 5/20

```
60000/60000 [=====] -
```

```
1s 11us/step - loss: 0.4877 - acc: 0.8751 - v  
al_loss: 0.4490 - val_acc: 0.8862  
Epoch 6/20  
60000/60000 [=====] -  
1s 11us/step - loss: 0.4617 - acc: 0.8800 - v  
al_loss: 0.4278 - val_acc: 0.8900  
Epoch 7/20  
60000/60000 [=====] -  
1s 10us/step - loss: 0.4425 - acc: 0.8838 - v  
al_loss: 0.4118 - val_acc: 0.8928  
Epoch 8/20  
60000/60000 [=====] -  
1s 11us/step - loss: 0.4275 - acc: 0.8860 - v  
al_loss: 0.3990 - val_acc: 0.8964  
Epoch 9/20  
60000/60000 [=====] -  
1s 11us/step - loss: 0.4155 - acc: 0.8887 - v  
al_loss: 0.3885 - val_acc: 0.8981  
Epoch 10/20  
60000/60000 [=====] -  
1s 10us/step - loss: 0.4055 - acc: 0.8906 - v  
al_loss: 0.3800 - val_acc: 0.9001  
Epoch 11/20  
60000/60000 [=====] -  
1s 11us/step - loss: 0.3971 - acc: 0.8926 - v  
al_loss: 0.3724 - val_acc: 0.9007  
Epoch 12/20  
60000/60000 [=====] -  
1s 11us/step - loss: 0.3898 - acc: 0.8940 - v  
al_loss: 0.3665 - val_acc: 0.9025  
Epoch 13/20  
60000/60000 [=====] -  
1s 10us/step - loss: 0.3834 - acc: 0.8952 - v  
al_loss: 0.3607 - val_acc: 0.9027  
Epoch 14/20  
60000/60000 [=====] -  
1s 10us/step - loss: 0.3779 - acc: 0.8966 - v
```

```
al_loss: 0.3560 - val_acc: 0.9039
Epoch 15/20
60000/60000 [=====] -
1s 11us/step - loss: 0.3729 - acc: 0.8978 - v
al_loss: 0.3517 - val_acc: 0.9050
Epoch 16/20
60000/60000 [=====] -
1s 11us/step - loss: 0.3684 - acc: 0.8985 - v
al_loss: 0.3474 - val_acc: 0.9058
Epoch 17/20
60000/60000 [=====] -
1s 11us/step - loss: 0.3643 - acc: 0.8997 - v
al_loss: 0.3440 - val_acc: 0.9066
Epoch 18/20
60000/60000 [=====] -
1s 11us/step - loss: 0.3606 - acc: 0.9005 - v
al_loss: 0.3405 - val_acc: 0.9073
Epoch 19/20
60000/60000 [=====] -
1s 11us/step - loss: 0.3572 - acc: 0.9011 - v
al_loss: 0.3374 - val_acc: 0.9084
Epoch 20/20
60000/60000 [=====] -
1s 11us/step - loss: 0.3540 - acc: 0.9019 - v
al_loss: 0.3350 - val_acc: 0.9094
```

In [32]:

```
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
```

```
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

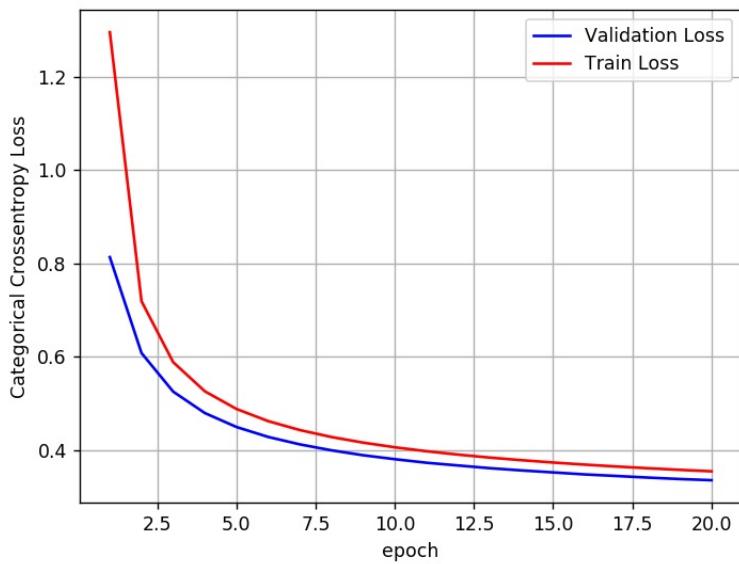
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.33496633795499803

Test accuracy: 0.9094



architecture 1 -hidden layers 2: input (28 *28) hidden layer-1 (512) layer-2 (128) output softmax 10

MLP + Sigmoid activation + SGDOptimizer

In [33]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape
Param #	
<hr/>	
<hr/>	
dense_2 (Dense)	(None, 512)
401920	
<hr/>	
<hr/>	
dense_3 (Dense)	(None, 128)
65664	
<hr/>	

```
dense_4 (Dense)           (None, 10)
    1290
=====
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
```

In [34]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

```
60000/60000 [=====] -
 3s 47us/step - loss: 2.2679 - acc: 0.2367 - val_loss: 2.2231 - val_acc: 0.3467
```

Epoch 2/20

```
60000/60000 [=====] -
 3s 46us/step - loss: 2.1810 - acc: 0.4511 - val_loss: 2.1272 - val_acc: 0.5366
```

Epoch 3/20

```
60000/60000 [=====] -
 3s 45us/step - loss: 2.0686 - acc: 0.5794 - val_loss: 1.9898 - val_acc: 0.6756
```

Epoch 4/20

```
60000/60000 [=====] -
 3s 45us/step - loss: 1.9073 - acc: 0.6430 - val_loss: 1.7961 - val_acc: 0.7165
```

```
Epoch 5/20
60000/60000 [=====] -
  3s 46us/step - loss: 1.6946 - acc: 0.6877 - val_loss: 1.5633 - val_acc: 0.7068
Epoch 6/20
60000/60000 [=====] -
  3s 46us/step - loss: 1.4628 - acc: 0.7181 - val_loss: 1.3351 - val_acc: 0.7570
Epoch 7/20
60000/60000 [=====] -
  3s 45us/step - loss: 1.2545 - acc: 0.7490 - val_loss: 1.1472 - val_acc: 0.7654
Epoch 8/20
60000/60000 [=====] -
  3s 48us/step - loss: 1.0885 - acc: 0.7708 - val_loss: 1.0024 - val_acc: 0.7828
Epoch 9/20
60000/60000 [=====] -
  3s 45us/step - loss: 0.9619 - acc: 0.7879 - val_loss: 0.8935 - val_acc: 0.8015
Epoch 10/20
60000/60000 [=====] -
  3s 48us/step - loss: 0.8652 - acc: 0.8024 - val_loss: 0.8078 - val_acc: 0.8176
Epoch 11/20
60000/60000 [=====] -
  3s 52us/step - loss: 0.7900 - acc: 0.8127 - val_loss: 0.7414 - val_acc: 0.8296
Epoch 12/20
60000/60000 [=====] -
  3s 51us/step - loss: 0.7299 - acc: 0.8229 - val_loss: 0.6879 - val_acc: 0.8385
Epoch 13/20
60000/60000 [=====] -
  3s 48us/step - loss: 0.6810 - acc: 0.8320 - val_loss: 0.6425 - val_acc: 0.8417
Epoch 14/20
```

```
60000/60000 [=====] -  
 3s 51us/step - loss: 0.6402 - acc: 0.8399 - v  
al_loss: 0.6053 - val_acc: 0.8486  
Epoch 15/20  
60000/60000 [=====] -  
 3s 51us/step - loss: 0.6058 - acc: 0.8465 - v  
al_loss: 0.5733 - val_acc: 0.8554  
Epoch 16/20  
60000/60000 [=====] -  
 3s 47us/step - loss: 0.5765 - acc: 0.8526 - v  
al_loss: 0.5461 - val_acc: 0.8604  
Epoch 17/20  
60000/60000 [=====] -  
 3s 49us/step - loss: 0.5512 - acc: 0.8577 - v  
al_loss: 0.5225 - val_acc: 0.8669  
Epoch 18/20  
60000/60000 [=====] -  
 3s 49us/step - loss: 0.5292 - acc: 0.8628 - v  
al_loss: 0.5016 - val_acc: 0.8715  
Epoch 19/20  
60000/60000 [=====] -  
 3s 51us/step - loss: 0.5099 - acc: 0.8667 - v  
al_loss: 0.4832 - val_acc: 0.8732  
Epoch 20/20  
60000/60000 [=====] -  
 3s 50us/step - loss: 0.4929 - acc: 0.8705 - v  
al_loss: 0.4670 - val_acc: 0.8783
```

In [35]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig,ax = plt.subplots(1,1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

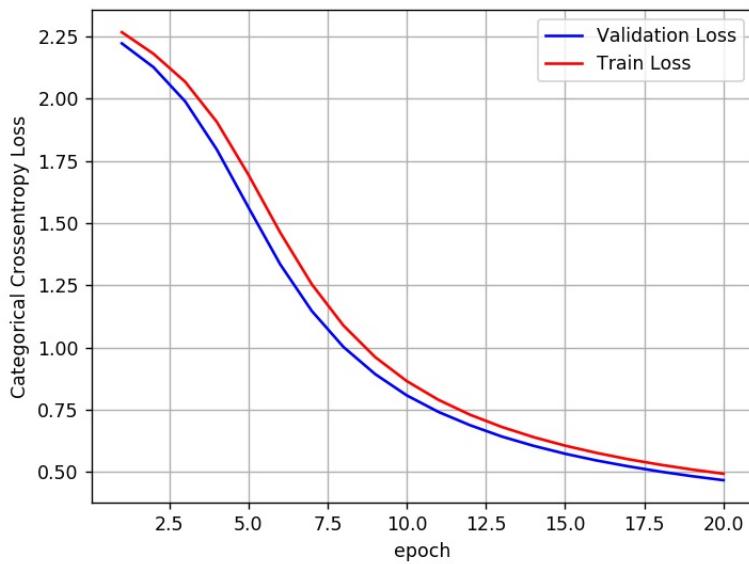
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.4670006361246109

Test accuracy: 0.8783



In [36]:

```
w_after = model_sigmoid.get_weights()

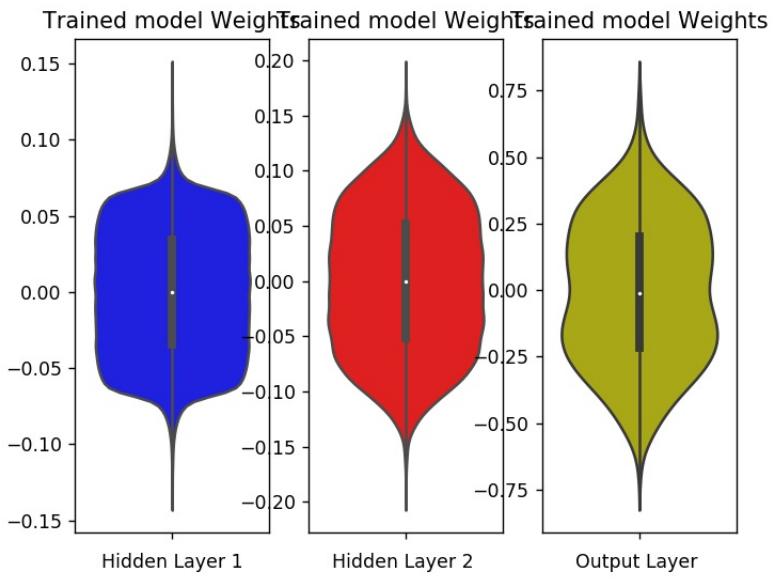
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + ADAM

In [37]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape
Param #	
dense_5 (Dense)	(None, 512)
	401920
dense_6 (Dense)	(None, 128)
	65664
dense_7 (Dense)	(None, 10)

1290

```
=====
=====
```

Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] -
4s 72us/step - loss: 0.5363 - acc: 0.8599 - val_loss: 0.2539 - val_acc: 0.9261

Epoch 2/20

60000/60000 [=====] -
4s 70us/step - loss: 0.2203 - acc: 0.9352 - val_loss: 0.1913 - val_acc: 0.9430

Epoch 3/20

60000/60000 [=====] -
5s 76us/step - loss: 0.1632 - acc: 0.9516 - val_loss: 0.1449 - val_acc: 0.9560

Epoch 4/20

60000/60000 [=====] -
4s 73us/step - loss: 0.1270 - acc: 0.9622 - val_loss: 0.1220 - val_acc: 0.9630

Epoch 5/20

60000/60000 [=====] -
5s 81us/step - loss: 0.1012 - acc: 0.9703 - val_loss: 0.1126 - val_acc: 0.9652

Epoch 6/20

60000/60000 [=====] -
4s 75us/step - loss: 0.0813 - acc: 0.9759 - val_loss: 0.0923 - val_acc: 0.9704

Epoch 7/20

60000/60000 [=====] -
4s 74us/step - loss: 0.0659 - acc: 0.9804 - val_loss: 0.0755 - val_acc: 0.9752

```
al_loss: 0.0855 - val_acc: 0.9737
Epoch 8/20
60000/60000 [=====] -
  5s 78us/step - loss: 0.0536 - acc: 0.9844 - v
al_loss: 0.0751 - val_acc: 0.9769
Epoch 9/20
60000/60000 [=====] -
  4s 74us/step - loss: 0.0434 - acc: 0.9874 - v
al_loss: 0.0702 - val_acc: 0.9786
Epoch 10/20
60000/60000 [=====] -
  5s 80us/step - loss: 0.0352 - acc: 0.9899 - v
al_loss: 0.0681 - val_acc: 0.9788
Epoch 11/20
60000/60000 [=====] -
  5s 75us/step - loss: 0.0281 - acc: 0.9927 - v
al_loss: 0.0635 - val_acc: 0.9800
Epoch 12/20
60000/60000 [=====] -
  5s 81us/step - loss: 0.0230 - acc: 0.9940 - v
al_loss: 0.0703 - val_acc: 0.9781
Epoch 13/20
60000/60000 [=====] -
  5s 75us/step - loss: 0.0178 - acc: 0.9954 - v
al_loss: 0.0759 - val_acc: 0.9771
Epoch 14/20
60000/60000 [=====] -
  5s 77us/step - loss: 0.0147 - acc: 0.9963 - v
al_loss: 0.0636 - val_acc: 0.9806
Epoch 15/20
60000/60000 [=====] -
  4s 75us/step - loss: 0.0105 - acc: 0.9977 - v
al_loss: 0.0671 - val_acc: 0.9802
Epoch 16/20
60000/60000 [=====] -
  4s 73us/step - loss: 0.0085 - acc: 0.9982 - v
al_loss: 0.0641 - val_acc: 0.9826
```

```
Epoch 17/20
60000/60000 [=====] -
  5s 76us/step - loss: 0.0078 - acc: 0.9983 - val_loss: 0.0732 - val_acc: 0.9798
Epoch 18/20
60000/60000 [=====] -
  4s 74us/step - loss: 0.0048 - acc: 0.9992 - val_loss: 0.0698 - val_acc: 0.9820
Epoch 19/20
60000/60000 [=====] -
  4s 70us/step - loss: 0.0046 - acc: 0.9990 - val_loss: 0.0676 - val_acc: 0.9827
Epoch 20/20
60000/60000 [=====] -
  4s 75us/step - loss: 0.0037 - acc: 0.9994 - val_loss: 0.0661 - val_acc: 0.9828
```

In [38]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the par
```

```

amter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

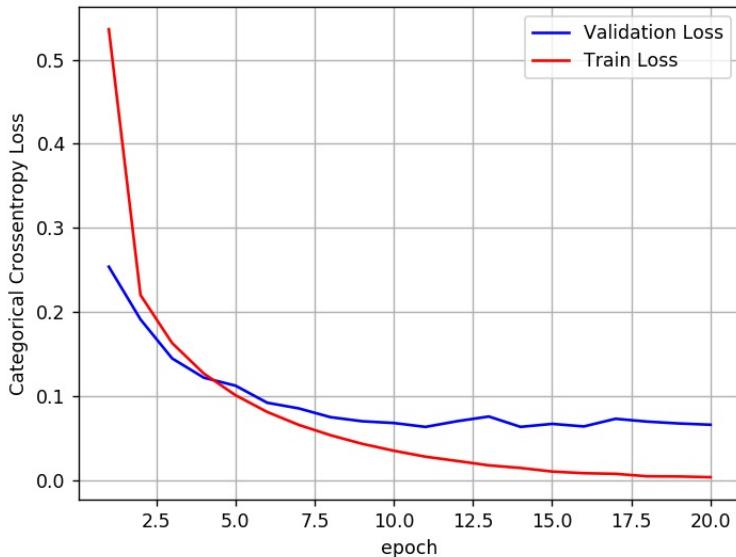
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06611316717128793

Test accuracy: 0.9828



In [39]:

```

w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)

```

```

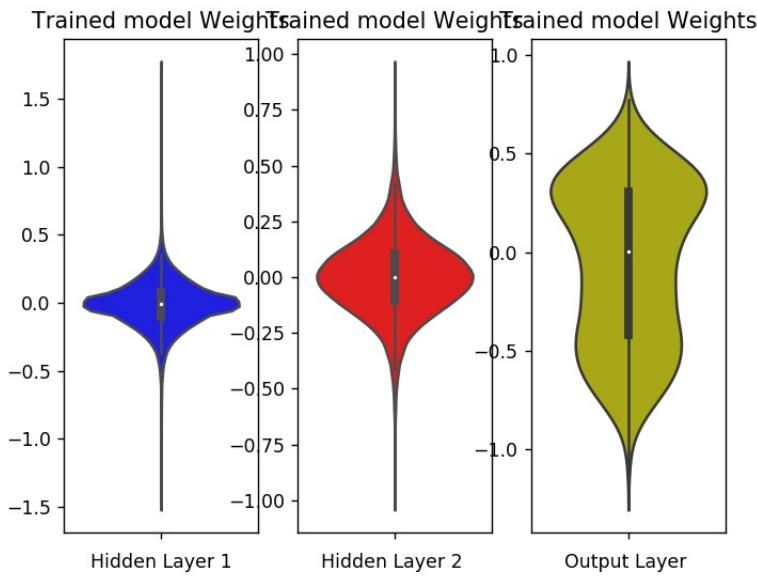
out_w = w_after[4].flatten().reshape(-1,1)

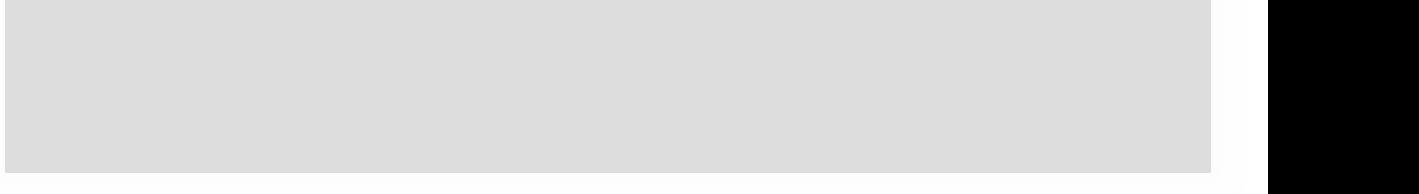
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```





MLP + ReLU +SGD

In [40]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(\theta, \sigma) = N(\theta, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(\theta, \sigma) = N(\theta, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape
Param #	
=====	=====
=====	=====
dense_8 (Dense)	(None, 512)
401920	

```
dense_9 (Dense)           (None, 128)
65664

dense_10 (Dense)          (None, 10)
1290
=====
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
```

In [41]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size,
epoches=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] -
3s 49us/step - loss: 0.7412 - acc: 0.7898 - val_loss: 0.3874 - val_acc: 0.8935
Epoch 2/20
60000/60000 [=====] -
3s 46us/step - loss: 0.3513 - acc: 0.9013 - val_loss: 0.3035 - val_acc: 0.9132
Epoch 3/20
60000/60000 [=====] -
3s 47us/step - loss: 0.2913 - acc: 0.9169 - val_loss: 0.2635 - val_acc: 0.9237
Epoch 4/20
```

```
60000/60000 [=====] -  
 3s 47us/step - loss: 0.2577 - acc: 0.9265 - v  
al_loss: 0.2407 - val_acc: 0.9316  
Epoch 5/20  
60000/60000 [=====] -  
 3s 46us/step - loss: 0.2345 - acc: 0.9335 - v  
al_loss: 0.2209 - val_acc: 0.9377  
Epoch 6/20  
60000/60000 [=====] -  
 3s 46us/step - loss: 0.2161 - acc: 0.9388 - v  
al_loss: 0.2079 - val_acc: 0.9421  
Epoch 7/20  
60000/60000 [=====] -  
 3s 47us/step - loss: 0.2012 - acc: 0.9433 - v  
al_loss: 0.1945 - val_acc: 0.9455  
Epoch 8/20  
60000/60000 [=====] -  
 3s 47us/step - loss: 0.1885 - acc: 0.9469 - v  
al_loss: 0.1870 - val_acc: 0.9475  
Epoch 9/20  
60000/60000 [=====] -  
 3s 46us/step - loss: 0.1775 - acc: 0.9502 - v  
al_loss: 0.1769 - val_acc: 0.9507  
Epoch 10/20  
60000/60000 [=====] -  
 3s 48us/step - loss: 0.1679 - acc: 0.9528 - v  
al_loss: 0.1686 - val_acc: 0.9519  
Epoch 11/20  
60000/60000 [=====] -  
 3s 51us/step - loss: 0.1594 - acc: 0.9552 - v  
al_loss: 0.1632 - val_acc: 0.9531  
Epoch 12/20  
60000/60000 [=====] -  
 3s 46us/step - loss: 0.1517 - acc: 0.9579 - v  
al_loss: 0.1561 - val_acc: 0.9551  
Epoch 13/20  
60000/60000 [=====] -
```

```
3s 52us/step - loss: 0.1448 - acc: 0.9595 - v
al_loss: 0.1510 - val_acc: 0.9565
Epoch 14/20
60000/60000 [=====] -
3s 47us/step - loss: 0.1384 - acc: 0.9616 - v
al_loss: 0.1460 - val_acc: 0.9567
Epoch 15/20
60000/60000 [=====] -
3s 52us/step - loss: 0.1327 - acc: 0.9630 - v
al_loss: 0.1412 - val_acc: 0.9580
Epoch 16/20
60000/60000 [=====] -
3s 52us/step - loss: 0.1274 - acc: 0.9646 - v
al_loss: 0.1374 - val_acc: 0.9597
Epoch 17/20
60000/60000 [=====] -
3s 47us/step - loss: 0.1224 - acc: 0.9660 - v
al_loss: 0.1340 - val_acc: 0.9613
Epoch 18/20
60000/60000 [=====] -
3s 49us/step - loss: 0.1179 - acc: 0.9675 - v
al_loss: 0.1302 - val_acc: 0.9617
Epoch 19/20
60000/60000 [=====] -
3s 56us/step - loss: 0.1137 - acc: 0.9689 - v
al_loss: 0.1255 - val_acc: 0.9625
Epoch 20/20
60000/60000 [=====] -
3s 50us/step - loss: 0.1097 - acc: 0.9699 - v
al_loss: 0.1244 - val_acc: 0.9630
```

In [42]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

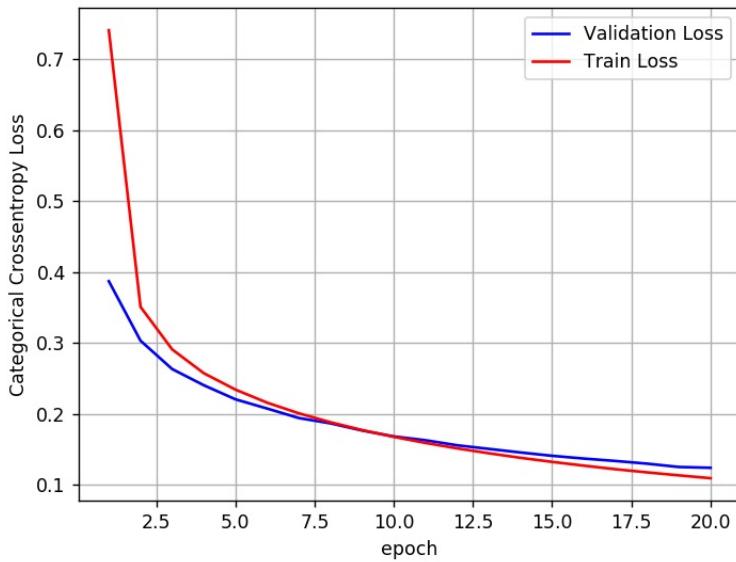
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12436589284352959

Test accuracy: 0.963



In [43]:

```
w_after = model_relu.get_weights()

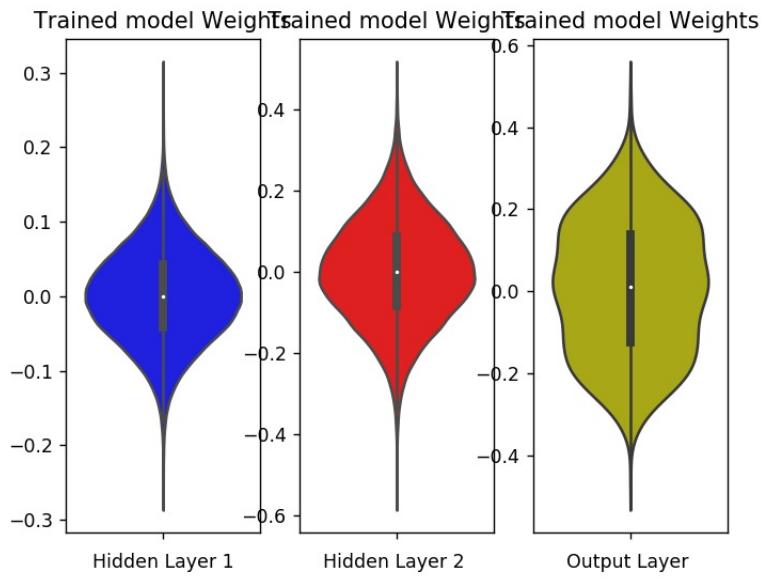
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + ADAM

In [44]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape
Param #	
dense_11 (Dense)	(None, 512)
	401920
dense_12 (Dense)	(None, 128)
	65664

```
dense_13 (Dense)           (None, 10)
    1290
=====
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
```

```
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] -
  4s 74us/step - loss: 0.2338 - acc: 0.9307 - val_loss: 0.1107 - val_acc: 0.9669
Epoch 2/20
60000/60000 [=====] -
  4s 68us/step - loss: 0.0876 - acc: 0.9735 - val_loss: 0.0840 - val_acc: 0.9726
Epoch 3/20
60000/60000 [=====] -
  4s 75us/step - loss: 0.0562 - acc: 0.9826 - val_loss: 0.0803 - val_acc: 0.9756
Epoch 4/20
60000/60000 [=====] -
  4s 73us/step - loss: 0.0385 - acc: 0.9880 - val_loss: 0.0762 - val_acc: 0.9771
Epoch 5/20
60000/60000 [=====] -
  4s 73us/step - loss: 0.0262 - acc: 0.9921 - val_loss: 0.0875 - val_acc: 0.9733
Epoch 6/20
60000/60000 [=====] -
  5s 75us/step - loss: 0.0199 - acc: 0.9936 - val_loss: 0.0745 - val_acc: 0.9771
```

Epoch 7/20
60000/60000 [=====] -
4s 73us/step - loss: 0.0179 - acc: 0.9942 - v
al_loss: 0.0926 - val_acc: 0.9744

Epoch 8/20
60000/60000 [=====] -
4s 74us/step - loss: 0.0158 - acc: 0.9949 - v
al_loss: 0.0781 - val_acc: 0.9798

Epoch 9/20
60000/60000 [=====] -
4s 74us/step - loss: 0.0131 - acc: 0.9955 - v
al_loss: 0.0934 - val_acc: 0.9766

Epoch 10/20
60000/60000 [=====] -
5s 80us/step - loss: 0.0106 - acc: 0.9966 - v
al_loss: 0.0751 - val_acc: 0.9809

Epoch 11/20
60000/60000 [=====] -
4s 74us/step - loss: 0.0117 - acc: 0.9962 - v
al_loss: 0.0722 - val_acc: 0.9811

Epoch 12/20
60000/60000 [=====] -
4s 72us/step - loss: 0.0085 - acc: 0.9973 - v
al_loss: 0.0893 - val_acc: 0.9799

Epoch 13/20
60000/60000 [=====] -
4s 74us/step - loss: 0.0117 - acc: 0.9963 - v
al_loss: 0.0819 - val_acc: 0.9800

Epoch 14/20
60000/60000 [=====] -
4s 72us/step - loss: 0.0103 - acc: 0.9967 - v
al_loss: 0.0917 - val_acc: 0.9795

Epoch 15/20
60000/60000 [=====] -
4s 72us/step - loss: 0.0099 - acc: 0.9968 - v
al_loss: 0.0864 - val_acc: 0.9808

Epoch 16/20

```
60000/60000 [=====] -  
 4s 72us/step - loss: 0.0056 - acc: 0.9981 - v  
al_loss: 0.0803 - val_acc: 0.9832  
Epoch 17/20  
60000/60000 [=====] -  
 4s 74us/step - loss: 0.0085 - acc: 0.9974 - v  
al_loss: 0.1044 - val_acc: 0.9797  
Epoch 18/20  
60000/60000 [=====] -  
 4s 74us/step - loss: 0.0095 - acc: 0.9970 - v  
al_loss: 0.0965 - val_acc: 0.9797  
Epoch 19/20  
60000/60000 [=====] -  
 4s 72us/step - loss: 0.0059 - acc: 0.9982 - v  
al_loss: 0.1049 - val_acc: 0.9802  
Epoch 20/20  
60000/60000 [=====] -  
 4s 74us/step - loss: 0.0082 - acc: 0.9971 - v  
al_loss: 0.0980 - val_acc: 0.9810
```

In [45]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig,ax = plt.subplots(1,1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')  
  
# list of epoch numbers  
x = list(range(1,nb_epoch+1))  
  
# print(history.history.keys())  
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])  
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y
```

```

        _test))

# we will get val_loss and val_acc only when you pass the par
amter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

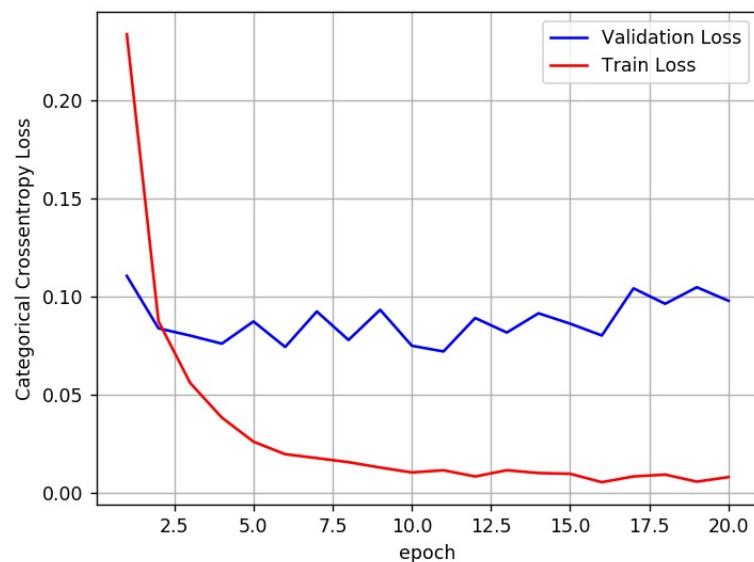
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of le
ngth equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09800619804806311

Test accuracy: 0.981



In [46]:

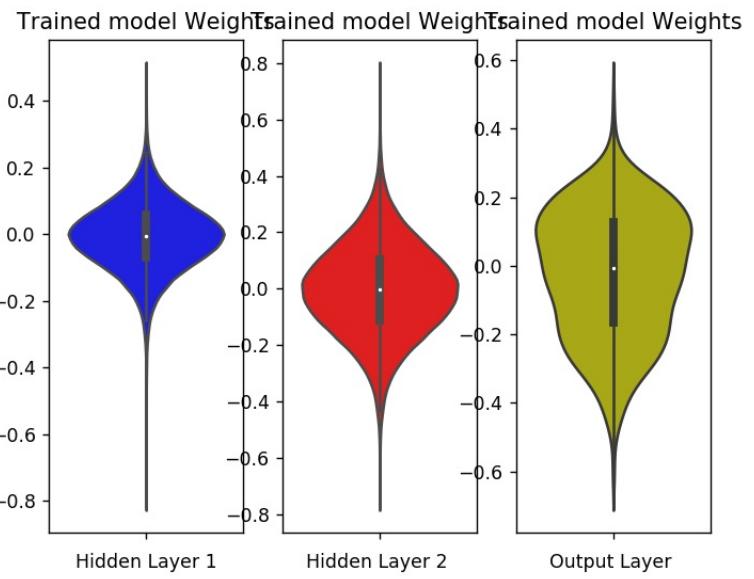
```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [47]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(0, 0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape
Param #	
=====	=====
dense_14 (Dense)	(None, 512)
401920	
=====	=====
batch_normalization_1 (Batch)	(None, 512)
2048	
=====	=====
dense_15 (Dense)	(None, 128)
65664	
=====	=====
batch_normalization_2 (Batch)	(None, 128)
512	
=====	=====
dense_16 (Dense)	(None, 10)
1290	
=====	=====
Total params:	471,434
Trainable params:	470,154
Non-trainable params:	1,280

In [48]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
                    metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_
```

```
size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_
test))
```

```
Train on 60000 samples, validate on 10000 samp
les
Epoch 1/20
60000/60000 [=====] -
  5s 89us/step - loss: 0.3067 - acc: 0.9107 - v
al_loss: 0.2121 - val_acc: 0.9393
Epoch 2/20
60000/60000 [=====] -
  5s 78us/step - loss: 0.1769 - acc: 0.9475 - v
al_loss: 0.1780 - val_acc: 0.9471
Epoch 3/20
60000/60000 [=====] -
  5s 78us/step - loss: 0.1389 - acc: 0.9593 - v
al_loss: 0.1526 - val_acc: 0.9537
Epoch 4/20
60000/60000 [=====] -
  5s 85us/step - loss: 0.1136 - acc: 0.9657 - v
al_loss: 0.1324 - val_acc: 0.9591
Epoch 5/20
60000/60000 [=====] -
  5s 82us/step - loss: 0.0973 - acc: 0.9708 - v
al_loss: 0.1286 - val_acc: 0.9611
Epoch 6/20
60000/60000 [=====] -
  5s 83us/step - loss: 0.0809 - acc: 0.9749 - v
al_loss: 0.1173 - val_acc: 0.9647
Epoch 7/20
60000/60000 [=====] -
  5s 85us/step - loss: 0.0675 - acc: 0.9791 - v
al_loss: 0.1133 - val_acc: 0.9642
Epoch 8/20
60000/60000 [=====] -
  5s 83us/step - loss: 0.0582 - acc: 0.9814 - v
al_loss: 0.1096 - val_acc: 0.9672
```

```
Epoch 9/20
60000/60000 [=====] -
  5s 86us/step - loss: 0.0516 - acc: 0.9836 - val_loss: 0.1088 - val_acc: 0.9674
Epoch 10/20
60000/60000 [=====] -
  5s 87us/step - loss: 0.0436 - acc: 0.9866 - val_loss: 0.1000 - val_acc: 0.9695
Epoch 11/20
60000/60000 [=====] -
  5s 87us/step - loss: 0.0417 - acc: 0.9863 - val_loss: 0.1067 - val_acc: 0.9682
Epoch 12/20
60000/60000 [=====] -
  5s 87us/step - loss: 0.0321 - acc: 0.9899 - val_loss: 0.0923 - val_acc: 0.9734
Epoch 13/20
60000/60000 [=====] -
  5s 90us/step - loss: 0.0301 - acc: 0.9900 - val_loss: 0.0995 - val_acc: 0.9714
Epoch 14/20
60000/60000 [=====] -
  5s 89us/step - loss: 0.0284 - acc: 0.9907 - val_loss: 0.0989 - val_acc: 0.9724
Epoch 15/20
60000/60000 [=====] -
  5s 88us/step - loss: 0.0262 - acc: 0.9913 - val_loss: 0.0993 - val_acc: 0.9716
Epoch 16/20
60000/60000 [=====] -
  5s 86us/step - loss: 0.0243 - acc: 0.9920 - val_loss: 0.1011 - val_acc: 0.9725
Epoch 17/20
60000/60000 [=====] -
  5s 82us/step - loss: 0.0194 - acc: 0.9938 - val_loss: 0.1056 - val_acc: 0.9716
Epoch 18/20
```

```
60000/60000 [=====] -  
5s 81us/step - loss: 0.0198 - acc: 0.9933 - v  
al_loss: 0.0990 - val_acc: 0.9751  
Epoch 19/20  
60000/60000 [=====] -  
5s 82us/step - loss: 0.0202 - acc: 0.9930 - v  
al_loss: 0.0937 - val_acc: 0.9756  
Epoch 20/20  
60000/60000 [=====] -  
5s 84us/step - loss: 0.0154 - acc: 0.9951 - v  
al_loss: 0.0997 - val_acc: 0.9741
```

In [49]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig,ax = plt.subplots(1,1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')  
  
# list of epoch numbers  
x = list(range(1,nb_epoch+1))  
  
# print(history.history.keys())  
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])  
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))  
  
# we will get val_loss and val_acc only when you pass the parameter validation_data  
# val_loss : validation loss  
# val_acc : validation accuracy  
  
# loss : training loss
```

```

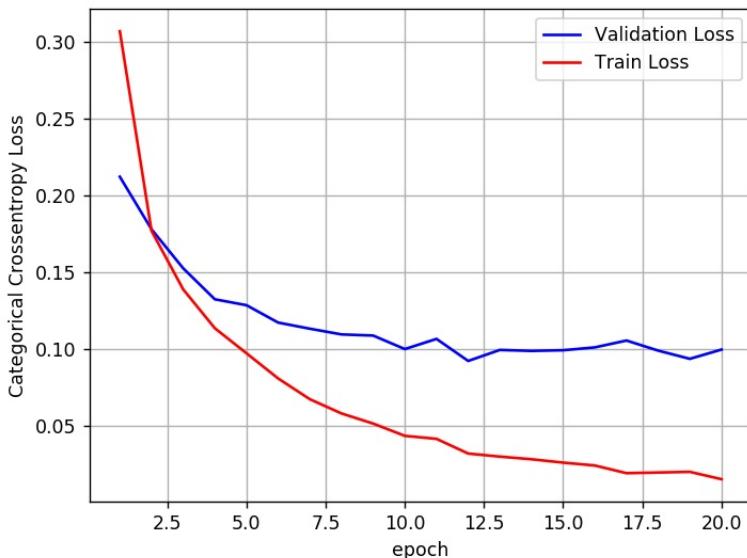
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09972703698020195

Test accuracy: 0.9741



In [50]:

```

w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")

```

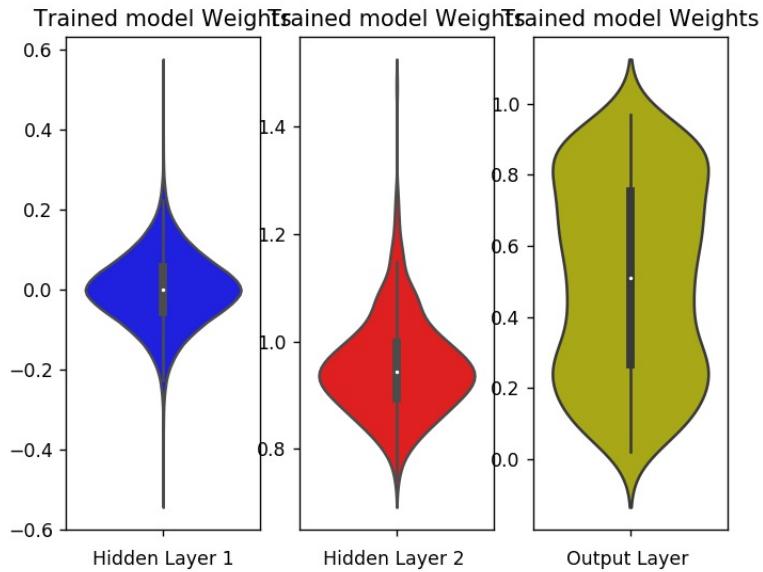
```

plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



5. MLP + Dropout + AdamOptimizer

In [51]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

```
WARNING:tensorflow:From C:\Users\Public\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version
```

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape
Param #	
dense_17 (Dense)	(None, 512)
	401920
batch_normalization_3 (Batch)	(None, 512)
	2048
dropout_1 (Dropout)	(None, 512)
	0
dense_18 (Dense)	(None, 128)
	65664
batch_normalization_4 (Batch)	(None, 128)
	512
dropout_2 (Dropout)	(None, 128)
	0
dense_19 (Dense)	(None, 10)
	1290

Total params: 471,434

```
Trainable params: 470,154  
Non-trainable params: 1,280
```

In [52]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/20
```

```
60000/60000 [=====] -  
6s 104us/step - loss: 0.6584 - acc: 0.7964 -  
val_loss: 0.2836 - val_acc: 0.9142
```

```
Epoch 2/20
```

```
60000/60000 [=====] -  
5s 85us/step - loss: 0.4215 - acc: 0.8706 - v  
al_loss: 0.2554 - val_acc: 0.9244
```

```
Epoch 3/20
```

```
60000/60000 [=====] -  
6s 93us/step - loss: 0.3776 - acc: 0.8845 - v  
al_loss: 0.2292 - val_acc: 0.9334
```

```
Epoch 4/20
```

```
60000/60000 [=====] -  
6s 95us/step - loss: 0.3514 - acc: 0.8932 - v  
al_loss: 0.2255 - val_acc: 0.9324
```

```
Epoch 5/20
```

```
60000/60000 [=====] -  
5s 90us/step - loss: 0.3320 - acc: 0.8992 - v  
al_loss: 0.2051 - val_acc: 0.9381
```

```
Epoch 6/20
```

```
60000/60000 [=====] -
```

```
6s 92us/step - loss: 0.3184 - acc: 0.9033 - v  
al_loss: 0.1978 - val_acc: 0.9413  
Epoch 7/20  
60000/60000 [=====] -  
6s 101us/step - loss: 0.3010 - acc: 0.9088 -  
val_loss: 0.1902 - val_acc: 0.9445  
Epoch 8/20  
60000/60000 [=====] -  
6s 97us/step - loss: 0.2923 - acc: 0.9117 - v  
al_loss: 0.1782 - val_acc: 0.9487  
Epoch 9/20  
60000/60000 [=====] -  
6s 98us/step - loss: 0.2766 - acc: 0.9165 - v  
al_loss: 0.1774 - val_acc: 0.9470  
Epoch 10/20  
60000/60000 [=====] -  
6s 99us/step - loss: 0.2645 - acc: 0.9198 - v  
al_loss: 0.1661 - val_acc: 0.9503  
Epoch 11/20  
60000/60000 [=====] -  
6s 96us/step - loss: 0.2557 - acc: 0.9233 - v  
al_loss: 0.1586 - val_acc: 0.9533  
Epoch 12/20  
60000/60000 [=====] -  
6s 97us/step - loss: 0.2459 - acc: 0.9266 - v  
al_loss: 0.1542 - val_acc: 0.9555  
Epoch 13/20  
60000/60000 [=====] -  
6s 98us/step - loss: 0.2326 - acc: 0.9286 - v  
al_loss: 0.1481 - val_acc: 0.9580  
Epoch 14/20  
60000/60000 [=====] -  
6s 96us/step - loss: 0.2265 - acc: 0.9311 - v  
al_loss: 0.1387 - val_acc: 0.9576  
Epoch 15/20  
60000/60000 [=====] -  
6s 99us/step - loss: 0.2204 - acc: 0.9329 - v
```

```
al_loss: 0.1299 - val_acc: 0.9595
Epoch 16/20
60000/60000 [=====] -
6s 92us/step - loss: 0.2067 - acc: 0.9388 - v
al_loss: 0.1251 - val_acc: 0.9629
Epoch 17/20
60000/60000 [=====] -
5s 89us/step - loss: 0.1953 - acc: 0.9405 - v
al_loss: 0.1180 - val_acc: 0.9635
Epoch 18/20
60000/60000 [=====] -
5s 92us/step - loss: 0.1912 - acc: 0.9416 - v
al_loss: 0.1169 - val_acc: 0.9648
Epoch 19/20
60000/60000 [=====] -
6s 92us/step - loss: 0.1825 - acc: 0.9441 - v
al_loss: 0.1127 - val_acc: 0.9665
Epoch 20/20
60000/60000 [=====] -
6s 95us/step - loss: 0.1753 - acc: 0.9470 - v
al_loss: 0.1099 - val_acc: 0.9679
```

In [53]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

```

# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
#                           epochs=nb_epoch, verbose=1, validation_data=(X_test, Y
#                           _test))

# we will get val_loss and val_acc only when you pass the par
# amter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

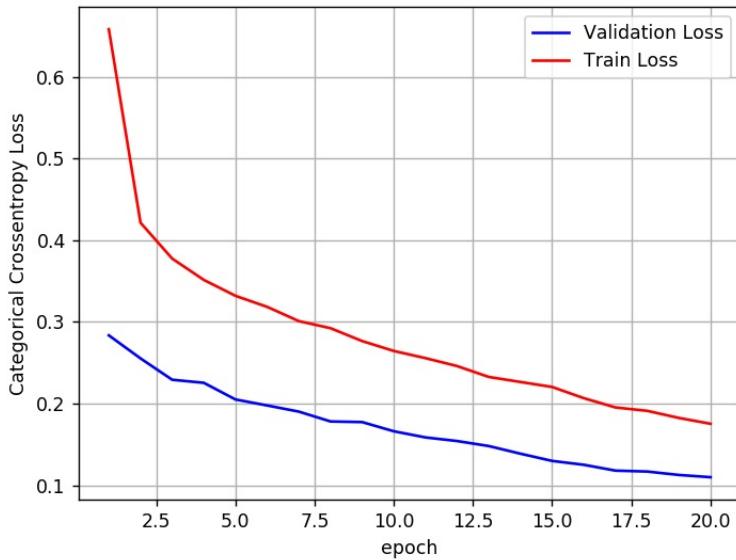
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of le
# ngth equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10993219921141863

Test accuracy: 0.9679



In [54]:

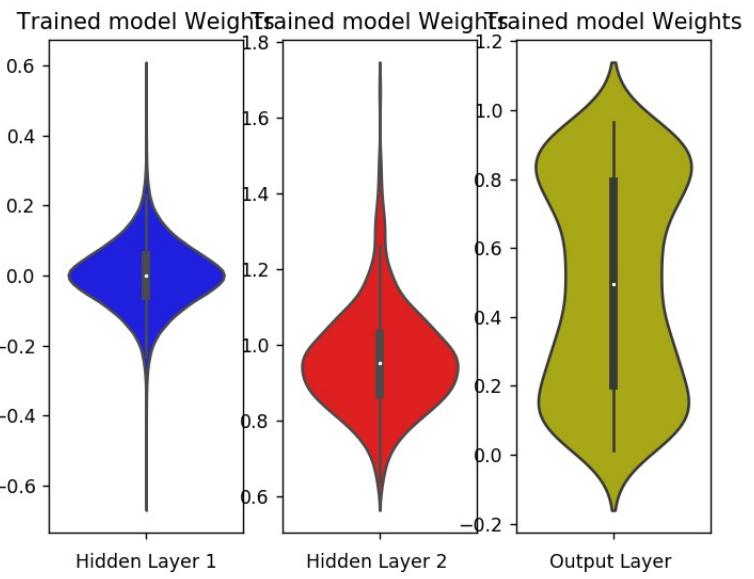
```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [55]:

```
from keras.optimizers import Adam, RMSprop, SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

In [56]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid', 'relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
```

```
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid,
n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
C:\Users\Public\Anaconda3\lib\site-packages\sk
learn\model_selection\_split.py:2053: FutureWa
rning: You should specify a value for 'cv' ins
tead of relying on the default value. The defa
ult value will change from 3 to 5 in version 0
.22.

warnings.warn(CV_WARNING, FutureWarning)
```

In [57]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_re
sult.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.974717 using {'activ': 'sigmoid'}
0.974717 (0.001036) with: {'activ': 'sigmoid'}
0.974400 (0.003102) with: {'activ': 'relu'}
```

**architecture 2 -hidden layers 3: input
(28 *28) hidden layer-1 (512) layer-2
(256) hidden layer-3 (64)output
softmax 10**

MLP + Sigmoid activation + SGDOptimizer

In [58]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(256, activation='sigmoid'))
model_sigmoid.add(Dense(64, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape
Param #	
=====	=====
=====	=====
dense_41 (Dense)	(None, 512)
401920	
=====	=====
=====	=====

```
dense_42 (Dense)           (None, 256)
131328

dense_43 (Dense)           (None, 64)
16448

dense_44 (Dense)           (None, 10)
650
=====
=====
Total params: 550,346
Trainable params: 550,346
Non-trainable params: 0
```

In [59]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 65us/step - loss: 2.3097 - acc: 0.1138 - val_loss: 2.2964 - val_acc: 0.1562
Epoch 2/20
60000/60000 [=====] - 3s 54us/step - loss: 2.2939 - acc: 0.1208 - val_loss: 2.2905 - val_acc: 0.2605
Epoch 3/20
```

60000/60000 [=====] -
3s 56us/step - loss: 2.2880 - acc: 0.1296 - v
al_loss: 2.2837 - val_acc: 0.1135
Epoch 4/20
60000/60000 [=====] -
4s 61us/step - loss: 2.2812 - acc: 0.1480 - v
al_loss: 2.2765 - val_acc: 0.1489
Epoch 5/20
60000/60000 [=====] -
4s 69us/step - loss: 2.2732 - acc: 0.1726 - v
al_loss: 2.2673 - val_acc: 0.2170
Epoch 6/20
60000/60000 [=====] -
4s 61us/step - loss: 2.2631 - acc: 0.2102 - v
al_loss: 2.2558 - val_acc: 0.3740
Epoch 7/20
60000/60000 [=====] -
4s 60us/step - loss: 2.2502 - acc: 0.2680 - v
al_loss: 2.2403 - val_acc: 0.2942
Epoch 8/20
60000/60000 [=====] -
3s 56us/step - loss: 2.2327 - acc: 0.3065 - v
al_loss: 2.2194 - val_acc: 0.3374
Epoch 9/20
60000/60000 [=====] -
4s 59us/step - loss: 2.2078 - acc: 0.3329 - v
al_loss: 2.1891 - val_acc: 0.3869
Epoch 10/20
60000/60000 [=====] -
4s 65us/step - loss: 2.1714 - acc: 0.3648 - v
al_loss: 2.1438 - val_acc: 0.3916
Epoch 11/20
60000/60000 [=====] -
4s 64us/step - loss: 2.1167 - acc: 0.3920 - v
al_loss: 2.0762 - val_acc: 0.3829
Epoch 12/20
60000/60000 [=====] -

```
4s 59us/step - loss: 2.0376 - acc: 0.4092 - val_loss: 1.9818 - val_acc: 0.4213
Epoch 13/20
60000/60000 [=====] -
4s 65us/step - loss: 1.9348 - acc: 0.4435 - val_loss: 1.8689 - val_acc: 0.4481
Epoch 14/20
60000/60000 [=====] -
4s 62us/step - loss: 1.8203 - acc: 0.4743 - val_loss: 1.7515 - val_acc: 0.4965
Epoch 15/20
60000/60000 [=====] -
3s 57us/step - loss: 1.7049 - acc: 0.5066 - val_loss: 1.6371 - val_acc: 0.5206
Epoch 16/20
60000/60000 [=====] -
4s 67us/step - loss: 1.5920 - acc: 0.5352 - val_loss: 1.5256 - val_acc: 0.5521
Epoch 17/20
60000/60000 [=====] -
4s 65us/step - loss: 1.4842 - acc: 0.5632 - val_loss: 1.4221 - val_acc: 0.5870
Epoch 18/20
60000/60000 [=====] -
4s 59us/step - loss: 1.3856 - acc: 0.5910 - val_loss: 1.3288 - val_acc: 0.6094
Epoch 19/20
60000/60000 [=====] -
4s 64us/step - loss: 1.2964 - acc: 0.6177 - val_loss: 1.2446 - val_acc: 0.6312
Epoch 20/20
60000/60000 [=====] -
4s 62us/step - loss: 1.2135 - acc: 0.6421 - val_loss: 1.1630 - val_acc: 0.6652
```

In [60]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

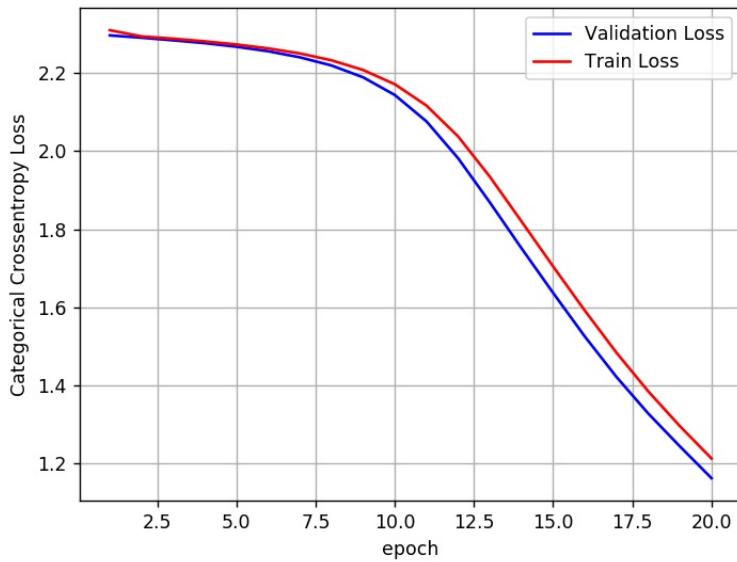
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 1.1629849285125733

Test accuracy: 0.6652



In [61]:

```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

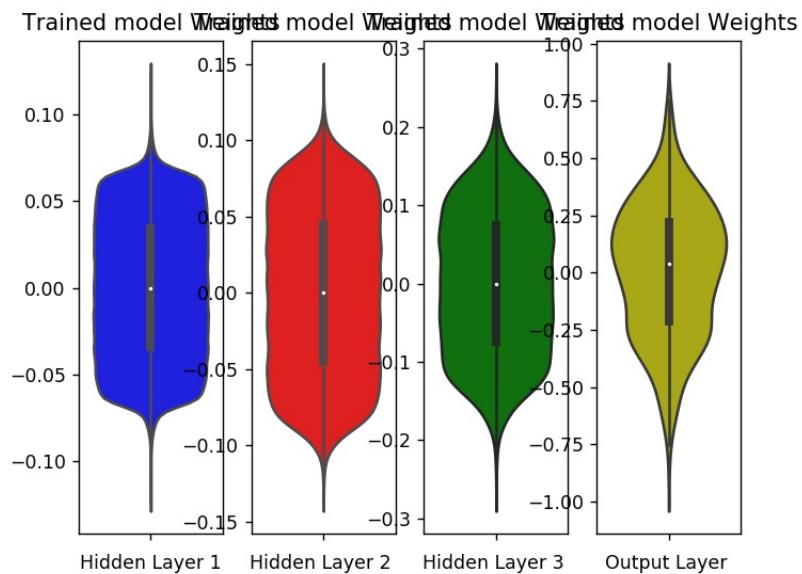
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()
```



MLP + Sigmoid activation + ADAM

In [62]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(256, activation='sigmoid'))
model_sigmoid.add(Dense(64, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape
Param #	
=====	=====
dense_45 (Dense)	(None, 512)
401920	
=====	=====
dense_46 (Dense)	(None, 256)
131328	
=====	=====

```
dense_47 (Dense)           (None, 64)
16448

dense_48 (Dense)           (None, 10)
650
=====
=====
Total params: 550,346
Trainable params: 550,346
Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] -
  6s 97us/step - loss: 0.6544 - acc: 0.8296 - val_loss: 0.2533 - val_acc: 0.9288
Epoch 2/20
60000/60000 [=====] -
  5s 83us/step - loss: 0.2156 - acc: 0.9372 - val_loss: 0.1816 - val_acc: 0.9445
Epoch 3/20
60000/60000 [=====] -
  5s 83us/step - loss: 0.1499 - acc: 0.9562 - val_loss: 0.1380 - val_acc: 0.9586
Epoch 4/20
60000/60000 [=====] -
  5s 86us/step - loss: 0.1127 - acc: 0.9674 - val_loss: 0.1036 - val_acc: 0.9675
Epoch 5/20
60000/60000 [=====] -
  5s 84us/step - loss: 0.0870 - acc: 0.9746 - val_loss: 0.0955 - val_acc: 0.9708
Epoch 6/20
60000/60000 [=====] -
```

```
5s 88us/step - loss: 0.0720 - acc: 0.9790 - v  
al_loss: 0.0856 - val_acc: 0.9748  
Epoch 7/20  
60000/60000 [=====] -  
5s 88us/step - loss: 0.0566 - acc: 0.9832 - v  
al_loss: 0.1005 - val_acc: 0.9678  
Epoch 8/20  
60000/60000 [=====] -  
5s 86us/step - loss: 0.0452 - acc: 0.9871 - v  
al_loss: 0.0793 - val_acc: 0.9754  
Epoch 9/20  
60000/60000 [=====] -  
5s 85us/step - loss: 0.0366 - acc: 0.9893 - v  
al_loss: 0.0803 - val_acc: 0.9760  
Epoch 10/20  
60000/60000 [=====] -  
5s 85us/step - loss: 0.0304 - acc: 0.9910 - v  
al_loss: 0.0658 - val_acc: 0.9798  
Epoch 11/20  
60000/60000 [=====] -  
5s 84us/step - loss: 0.0246 - acc: 0.9929 - v  
al_loss: 0.0756 - val_acc: 0.9775  
Epoch 12/20  
60000/60000 [=====] -  
5s 90us/step - loss: 0.0201 - acc: 0.9945 - v  
al_loss: 0.0799 - val_acc: 0.9773  
Epoch 13/20  
60000/60000 [=====] -  
5s 89us/step - loss: 0.0159 - acc: 0.9959 - v  
al_loss: 0.0721 - val_acc: 0.9803  
Epoch 14/20  
60000/60000 [=====] -  
5s 87us/step - loss: 0.0135 - acc: 0.9963 - v  
al_loss: 0.0853 - val_acc: 0.9769  
Epoch 15/20  
60000/60000 [=====] -  
5s 87us/step - loss: 0.0131 - acc: 0.9962 - v
```

```
al_loss: 0.0786 - val_acc: 0.9783
Epoch 16/20
60000/60000 [=====] -
 5s 84us/step - loss: 0.0096 - acc: 0.9976 - v
al_loss: 0.0774 - val_acc: 0.9785
Epoch 17/20
60000/60000 [=====] -
 5s 84us/step - loss: 0.0098 - acc: 0.9972 - v
al_loss: 0.0729 - val_acc: 0.9802
Epoch 18/20
60000/60000 [=====] -
 5s 86us/step - loss: 0.0087 - acc: 0.9975 - v
al_loss: 0.0866 - val_acc: 0.9783
Epoch 19/20
60000/60000 [=====] -
 5s 84us/step - loss: 0.0072 - acc: 0.9981 - v
al_loss: 0.0711 - val_acc: 0.9815
Epoch 20/20
60000/60000 [=====] -
 5s 87us/step - loss: 0.0064 - acc: 0.9983 - v
al_loss: 0.0770 - val_acc: 0.9812
```

In [63]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

```

# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
#                           epochs=nb_epoch, verbose=1, validation_data=(X_test, Y
#                           _test))

# we will get val_loss and val_acc only when you pass the par
# amter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

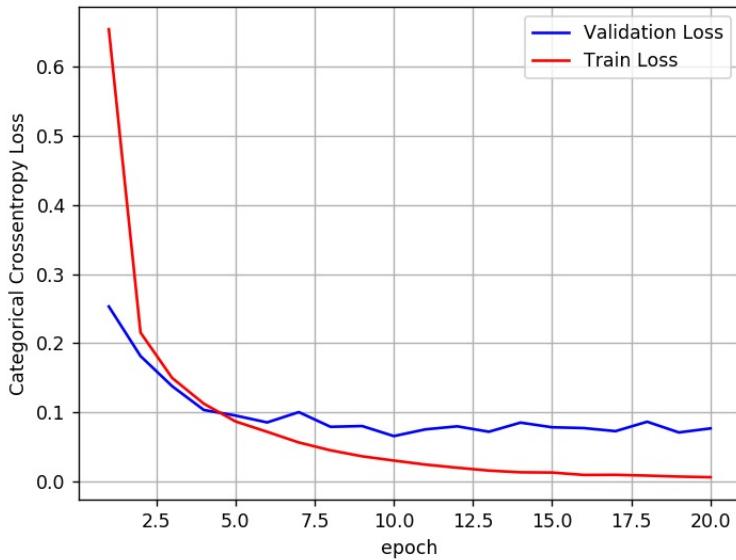
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of le
# ngth equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07696367663278943

Test accuracy: 0.9812



In [64]:

```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

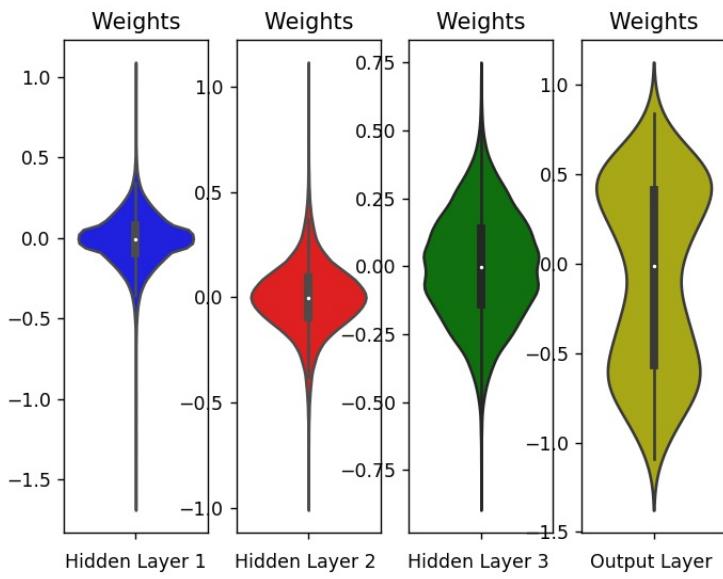
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU +SGD

In [65]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(\theta, \sigma) = N(\theta, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.088 \Rightarrow N(\theta, \sigma) = N(\theta, 0.088)$ 
# h3 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.176 \Rightarrow N(\theta, \sigma) = N(\theta, 0.176)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.164 \Rightarrow N(\theta, \sigma) = N(\theta, 0.164)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)) )
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)) )

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape
Param #	
=====	=====
=====	

```
dense_49 (Dense)           (None, 512)
401920

dense_50 (Dense)           (None, 256)
131328

dense_51 (Dense)           (None, 64)
16448

dense_52 (Dense)           (None, 10)
650
=====
=====
Total params: 550,346
Trainable params: 550,346
Non-trainable params: 0
```

In [66]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size,
epoches=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] -
5s 76us/step - loss: 0.5967 - acc: 0.8265 - val_loss: 0.3096 - val_acc: 0.9089
Epoch 2/20
```

60000/60000 [=====] -
4s 59us/step - loss: 0.2810 - acc: 0.9181 - v
al_loss: 0.2406 - val_acc: 0.9304

Epoch 3/20

60000/60000 [=====] -
4s 65us/step - loss: 0.2270 - acc: 0.9340 - v
al_loss: 0.2068 - val_acc: 0.9388

Epoch 4/20

60000/60000 [=====] -
4s 60us/step - loss: 0.1941 - acc: 0.9442 - v
al_loss: 0.1848 - val_acc: 0.9441

Epoch 5/20

60000/60000 [=====] -
3s 58us/step - loss: 0.1718 - acc: 0.9505 - v
al_loss: 0.1701 - val_acc: 0.9495

Epoch 6/20

60000/60000 [=====] -
4s 59us/step - loss: 0.1550 - acc: 0.9554 - v
al_loss: 0.1567 - val_acc: 0.9522

Epoch 7/20

60000/60000 [=====] -
4s 64us/step - loss: 0.1410 - acc: 0.9603 - v
al_loss: 0.1476 - val_acc: 0.9551

Epoch 8/20

60000/60000 [=====] -
4s 65us/step - loss: 0.1295 - acc: 0.9632 - v
al_loss: 0.1394 - val_acc: 0.9580

Epoch 9/20

60000/60000 [=====] -
4s 64us/step - loss: 0.1197 - acc: 0.9663 - v
al_loss: 0.1330 - val_acc: 0.9590

Epoch 10/20

60000/60000 [=====] -
4s 65us/step - loss: 0.1113 - acc: 0.9687 - v
al_loss: 0.1296 - val_acc: 0.9593

Epoch 11/20

60000/60000 [=====] -

```
4s 58us/step - loss: 0.1037 - acc: 0.9713 - v  
al_loss: 0.1212 - val_acc: 0.9612  
Epoch 12/20  
60000/60000 [=====] -  
4s 67us/step - loss: 0.0971 - acc: 0.9729 - v  
al_loss: 0.1180 - val_acc: 0.9627  
Epoch 13/20  
60000/60000 [=====] -  
4s 64us/step - loss: 0.0914 - acc: 0.9745 - v  
al_loss: 0.1145 - val_acc: 0.9637  
Epoch 14/20  
60000/60000 [=====] -  
4s 62us/step - loss: 0.0860 - acc: 0.9766 - v  
al_loss: 0.1121 - val_acc: 0.9655  
Epoch 15/20  
60000/60000 [=====] -  
4s 63us/step - loss: 0.0810 - acc: 0.9776 - v  
al_loss: 0.1100 - val_acc: 0.9646  
Epoch 16/20  
60000/60000 [=====] -  
4s 62us/step - loss: 0.0765 - acc: 0.9794 - v  
al_loss: 0.1091 - val_acc: 0.9649  
Epoch 17/20  
60000/60000 [=====] -  
4s 62us/step - loss: 0.0722 - acc: 0.9803 - v  
al_loss: 0.1055 - val_acc: 0.9665  
Epoch 18/20  
60000/60000 [=====] -  
4s 61us/step - loss: 0.0685 - acc: 0.9815 - v  
al_loss: 0.1018 - val_acc: 0.9668  
Epoch 19/20  
60000/60000 [=====] -  
4s 61us/step - loss: 0.0652 - acc: 0.9830 - v  
al_loss: 0.1023 - val_acc: 0.9676  
Epoch 20/20  
60000/60000 [=====] -  
4s 61us/step - loss: 0.0619 - acc: 0.9838 - v
```

```
al_loss: 0.1000 - val_acc: 0.9685
```

In [67]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

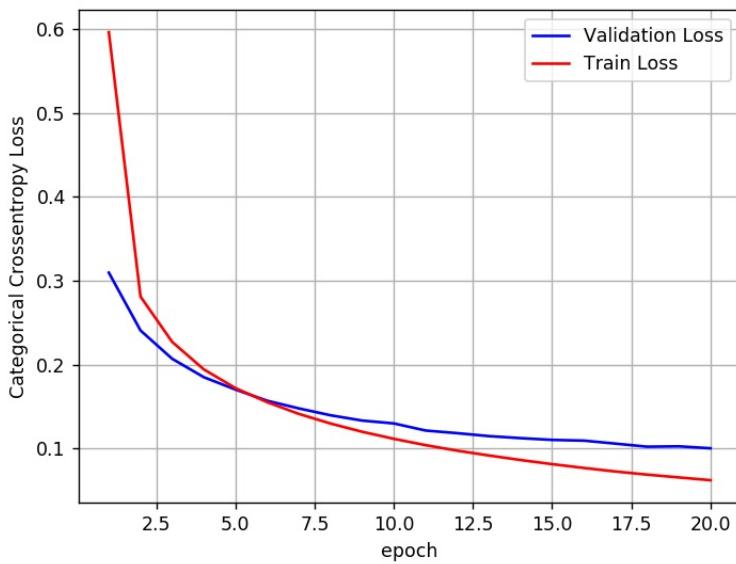
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09995875298418104

Test accuracy: 0.9685



In [68]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

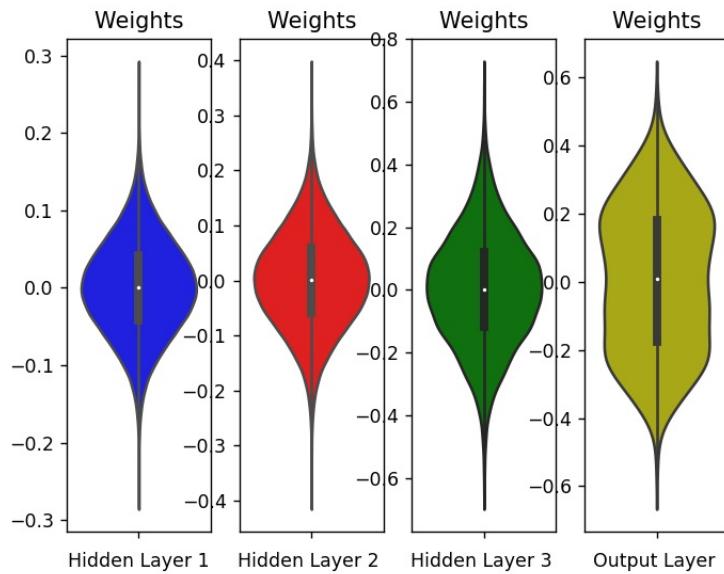
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()
```



MLP + ReLU + ADAM

In [69]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)) )
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape
Param #	
dense_53 (Dense)	(None, 512)
	401920
dense_54 (Dense)	(None, 256)

131328

dense_55 (Dense)	(None, 64)
16448	

dense_56 (Dense)	(None, 10)
650	

Total params: 550,346

Trainable params: 550,346

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] -
6s 97us/step - loss: 0.2275 - acc: 0.9318 - val_loss: 0.1154 - val_acc: 0.9637

Epoch 2/20

60000/60000 [=====] -
5s 89us/step - loss: 0.0838 - acc: 0.9738 - val_loss: 0.0851 - val_acc: 0.9739

Epoch 3/20

60000/60000 [=====] -
5s 84us/step - loss: 0.0511 - acc: 0.9840 - val_loss: 0.0930 - val_acc: 0.9708

Epoch 4/20

60000/60000 [=====] -
5s 86us/step - loss: 0.0390 - acc: 0.9878 - val_loss: 0.0712 - val_acc: 0.9768

Epoch 5/20

60000/60000 [=====] -

```
5s 85us/step - loss: 0.0286 - acc: 0.9902 - v  
al_loss: 0.0804 - val_acc: 0.9760  
Epoch 6/20  
60000/60000 [=====] -  
5s 88us/step - loss: 0.0233 - acc: 0.9923 - v  
al_loss: 0.0780 - val_acc: 0.9787  
Epoch 7/20  
60000/60000 [=====] -  
5s 88us/step - loss: 0.0206 - acc: 0.9935 - v  
al_loss: 0.0834 - val_acc: 0.9765  
Epoch 8/20  
60000/60000 [=====] -  
5s 88us/step - loss: 0.0176 - acc: 0.9941 - v  
al_loss: 0.0863 - val_acc: 0.9786  
Epoch 9/20  
60000/60000 [=====] -  
5s 84us/step - loss: 0.0154 - acc: 0.9949 - v  
al_loss: 0.0880 - val_acc: 0.9798  
Epoch 10/20  
60000/60000 [=====] -  
5s 84us/step - loss: 0.0182 - acc: 0.9938 - v  
al_loss: 0.0997 - val_acc: 0.9769  
Epoch 11/20  
60000/60000 [=====] -  
5s 85us/step - loss: 0.0126 - acc: 0.9960 - v  
al_loss: 0.0919 - val_acc: 0.9792  
Epoch 12/20  
60000/60000 [=====] -  
5s 85us/step - loss: 0.0120 - acc: 0.9958 - v  
al_loss: 0.0902 - val_acc: 0.9799  
Epoch 13/20  
60000/60000 [=====] -  
5s 83us/step - loss: 0.0159 - acc: 0.9947 - v  
al_loss: 0.0885 - val_acc: 0.9781  
Epoch 14/20  
60000/60000 [=====] -  
5s 86us/step - loss: 0.0129 - acc: 0.9958 - v
```

```
al_loss: 0.1049 - val_acc: 0.9775
Epoch 15/20
60000/60000 [=====] -
  5s 83us/step - loss: 0.0058 - acc: 0.9982 - v
al_loss: 0.1280 - val_acc: 0.9757
Epoch 16/20
60000/60000 [=====] -
  5s 87us/step - loss: 0.0142 - acc: 0.9956 - v
al_loss: 0.1147 - val_acc: 0.9766
Epoch 17/20
60000/60000 [=====] -
  5s 88us/step - loss: 0.0126 - acc: 0.9961 - v
al_loss: 0.0822 - val_acc: 0.9831
Epoch 18/20
60000/60000 [=====] -
  5s 84us/step - loss: 0.0060 - acc: 0.9980 - v
al_loss: 0.0982 - val_acc: 0.9799
Epoch 19/20
60000/60000 [=====] -
  5s 84us/step - loss: 0.0086 - acc: 0.9971 - v
al_loss: 0.1200 - val_acc: 0.9743
Epoch 20/20
60000/60000 [=====] -
  5s 85us/step - loss: 0.0140 - acc: 0.9957 - v
al_loss: 0.0907 - val_acc: 0.9807
```

In [70]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
```

```

x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

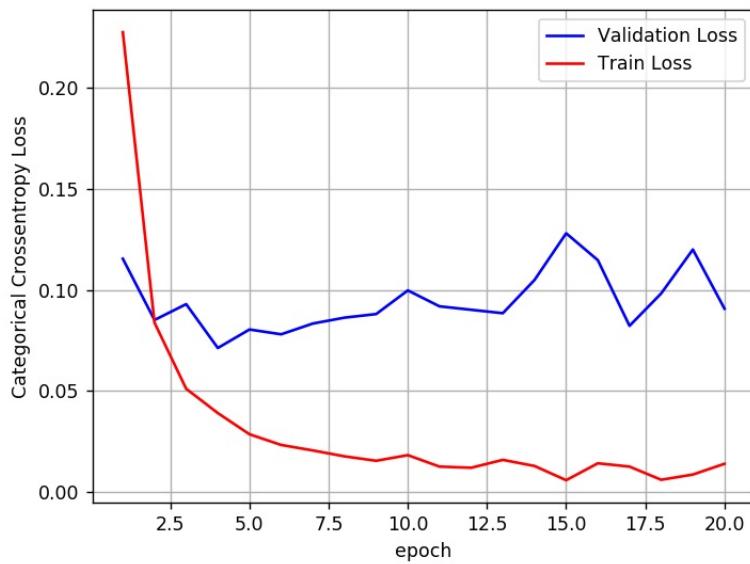
Test score: 0.09070437623546104

Test accuracy: 0.9807

```

C:\Users\Public\Anaconda3\lib\site-packages\matplotlib\pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).
    max_open_warning, RuntimeWarning)

```



In [71]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

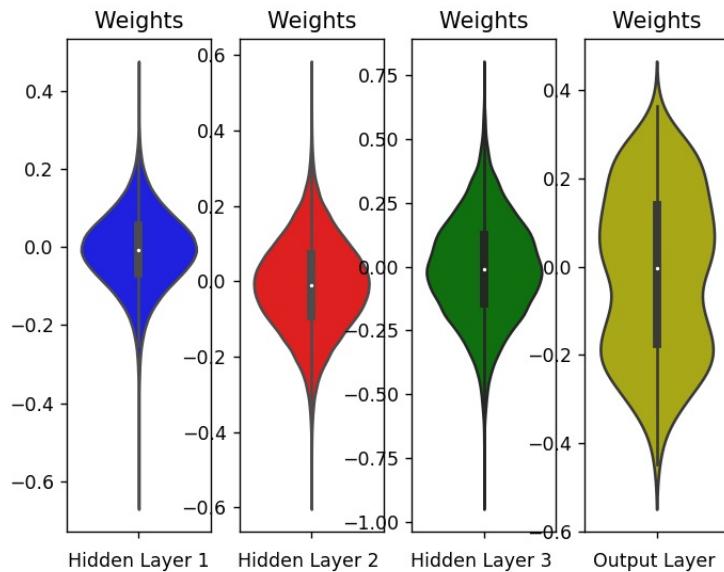
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [72]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.051 \Rightarrow N(\theta, \sigma) = N(0, 0.051)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.079 \Rightarrow N(\theta, \sigma) = N(0, 0.079)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.079, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))
```

```
model_batch.summary()
```

Layer (type)	Output Shape
Param #	
dense_57 (Dense)	(None, 512)
	401920
batch_normalization_5 (Batch)	(None, 512)
	2048
dense_58 (Dense)	(None, 256)
	131328
batch_normalization_6 (Batch)	(None, 256)
	1024
dense_59 (Dense)	(None, 64)
	16448
batch_normalization_7 (Batch)	(None, 64)
	256
dense_60 (Dense)	(None, 10)
	650

```
Total params: 553,674  
Trainable params: 552,010  
Non-trainable params: 1,664
```

In [73]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 130us/step - loss: 0.2636 - acc: 0.9242 -
val_loss: 0.1515 - val_acc: 0.9575
Epoch 2/20
60000/60000 [=====] - 6s 103us/step - loss: 0.1288 - acc: 0.9616 -
val_loss: 0.1129 - val_acc: 0.9672
Epoch 3/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0869 - acc: 0.9739 -
val_loss: 0.0858 - val_acc: 0.9748
Epoch 4/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0623 - acc: 0.9806 -
val_loss: 0.0942 - val_acc: 0.9704
Epoch 5/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0477 - acc: 0.9850 -
val_loss: 0.0814 - val_acc: 0.9749
Epoch 6/20
```

```
60000/60000 [=====] -  
 6s 100us/step - loss: 0.0365 - acc: 0.9885 -  
val_loss: 0.0719 - val_acc: 0.9779  
Epoch 7/20  
60000/60000 [=====] -  
 6s 99us/step - loss: 0.0361 - acc: 0.9884 - v  
al_loss: 0.0780 - val_acc: 0.9752  
Epoch 8/20  
60000/60000 [=====] -  
 6s 100us/step - loss: 0.0263 - acc: 0.9913 -  
val_loss: 0.0771 - val_acc: 0.9776  
Epoch 9/20  
60000/60000 [=====] -  
 6s 97us/step - loss: 0.0225 - acc: 0.9923 - v  
al_loss: 0.0805 - val_acc: 0.9780  
Epoch 10/20  
60000/60000 [=====] -  
 6s 106us/step - loss: 0.0217 - acc: 0.9927 -  
val_loss: 0.0893 - val_acc: 0.9748  
Epoch 11/20  
60000/60000 [=====] -  
 6s 99us/step - loss: 0.0199 - acc: 0.9931 - v  
al_loss: 0.0812 - val_acc: 0.9770  
Epoch 12/20  
60000/60000 [=====] -  
 6s 100us/step - loss: 0.0161 - acc: 0.9945 -  
val_loss: 0.0981 - val_acc: 0.9731  
Epoch 13/20  
60000/60000 [=====] -  
 6s 99us/step - loss: 0.0155 - acc: 0.9947 - v  
al_loss: 0.0855 - val_acc: 0.9760  
Epoch 14/20  
60000/60000 [=====] -  
 6s 99us/step - loss: 0.0157 - acc: 0.9948 - v  
al_loss: 0.0783 - val_acc: 0.9782  
Epoch 15/20  
60000/60000 [=====] -
```

```
6s 100us/step - loss: 0.0141 - acc: 0.9952 -
val_loss: 0.0789 - val_acc: 0.9797
Epoch 16/20
60000/60000 [=====] -
6s 98us/step - loss: 0.0116 - acc: 0.9961 - v
al_loss: 0.0850 - val_acc: 0.9781
Epoch 17/20
60000/60000 [=====] -
6s 99us/step - loss: 0.0131 - acc: 0.9957 - v
al_loss: 0.0817 - val_acc: 0.9802
Epoch 18/20
60000/60000 [=====] -
6s 107us/step - loss: 0.0122 - acc: 0.9957 -
val_loss: 0.0848 - val_acc: 0.9790
Epoch 19/20
60000/60000 [=====] -
6s 101us/step - loss: 0.0128 - acc: 0.9955 -
val_loss: 0.0880 - val_acc: 0.9794
Epoch 20/20
60000/60000 [=====] -
6s 107us/step - loss: 0.0122 - acc: 0.9957 -
val_loss: 0.0839 - val_acc: 0.9790
```

In [74]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
```

```

# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

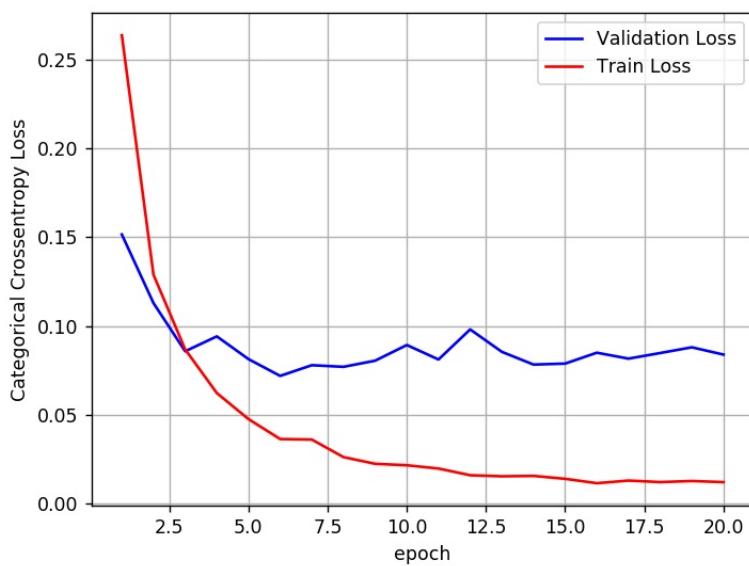
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08393007170943602

Test accuracy: 0.979



In [75]:

```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

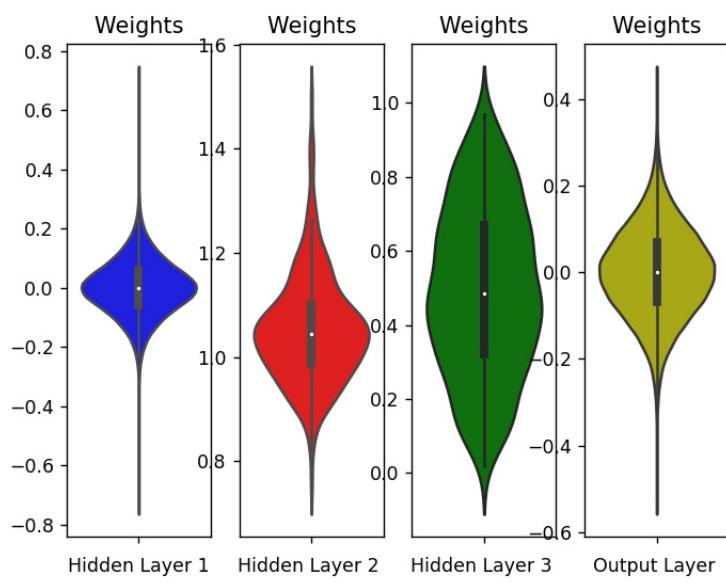
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5. MLP + Dropout + AdamOptimizer

In [76]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None) ))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.079, seed=None) ))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape
Param #	
dense_61 (Dense)	(None, 512)
	401920
batch_normalization_8 (Batch Normalization)	(None, 512)
	2048
dropout_3 (Dropout)	(None, 512)
	0
dense_62 (Dense)	(None, 256)
	131328
batch_normalization_9 (Batch Normalization)	(None, 256)
	1024
dropout_4 (Dropout)	(None, 256)
	0
dense_63 (Dense)	(None, 64)
	16448
batch_normalization_10 (Batch Normalization)	(None, 64)
	256
dropout_5 (Dropout)	(None, 64)

0

```
dense_64 (Dense)           (None, 10)
```

```
650
```

```
=====
```

```
=====
```

```
Total params: 553,674
```

```
Trainable params: 552,010
```

```
Non-trainable params: 1,664
```

In [77]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/20
```

```
60000/60000 [=====] -
```

```
9s 144us/step - loss: 0.5961 - acc: 0.8192 -
```

```
val_loss: 0.2633 - val_acc: 0.9214
```

```
Epoch 2/20
```

```
60000/60000 [=====] -
```

```
7s 116us/step - loss: 0.3927 - acc: 0.8853 -
```

```
val_loss: 0.2321 - val_acc: 0.9319
```

```
Epoch 3/20
```

```
60000/60000 [=====] -
```

```
7s 119us/step - loss: 0.3416 - acc: 0.9001 -
```

```
val_loss: 0.2031 - val_acc: 0.9399
```

```
Epoch 4/20
```

```
60000/60000 [=====] -
```

```
    7s 114us/step - loss: 0.3048 - acc: 0.9110 -
val_loss: 0.1796 - val_acc: 0.9491
Epoch 5/20
60000/60000 [=====] -
    7s 114us/step - loss: 0.2755 - acc: 0.9202 -
val_loss: 0.1636 - val_acc: 0.9520
Epoch 6/20
60000/60000 [=====] -
    7s 118us/step - loss: 0.2457 - acc: 0.9291 -
val_loss: 0.1442 - val_acc: 0.9571
Epoch 7/20
60000/60000 [=====] -
    7s 117us/step - loss: 0.2292 - acc: 0.9338 -
val_loss: 0.1306 - val_acc: 0.9610
Epoch 8/20
60000/60000 [=====] -
    7s 113us/step - loss: 0.2069 - acc: 0.9400 -
val_loss: 0.1226 - val_acc: 0.9638
Epoch 9/20
60000/60000 [=====] -
    7s 113us/step - loss: 0.1890 - acc: 0.9451 -
val_loss: 0.1113 - val_acc: 0.9681
Epoch 10/20
60000/60000 [=====] -
    7s 110us/step - loss: 0.1775 - acc: 0.9490 -
val_loss: 0.1005 - val_acc: 0.9714
Epoch 11/20
60000/60000 [=====] -
    7s 110us/step - loss: 0.1631 - acc: 0.9534 -
val_loss: 0.0971 - val_acc: 0.9720
Epoch 12/20
60000/60000 [=====] -
    7s 116us/step - loss: 0.1509 - acc: 0.9572 -
val_loss: 0.0903 - val_acc: 0.9743
Epoch 13/20
60000/60000 [=====] -
    7s 122us/step - loss: 0.1459 - acc: 0.9579 -
```

```
val_loss: 0.0881 - val_acc: 0.9754
Epoch 14/20
60000/60000 [=====] -
    7s 113us/step - loss: 0.1382 - acc: 0.9595 -
val_loss: 0.0819 - val_acc: 0.9762
Epoch 15/20
60000/60000 [=====] -
    7s 111us/step - loss: 0.1291 - acc: 0.9618 -
val_loss: 0.0786 - val_acc: 0.9776
Epoch 16/20
60000/60000 [=====] -
    7s 114us/step - loss: 0.1250 - acc: 0.9643 -
val_loss: 0.0712 - val_acc: 0.9796
Epoch 17/20
60000/60000 [=====] -
    7s 113us/step - loss: 0.1163 - acc: 0.9667 -
val_loss: 0.0745 - val_acc: 0.9791
Epoch 18/20
60000/60000 [=====] -
    7s 111us/step - loss: 0.1096 - acc: 0.9688 -
val_loss: 0.0721 - val_acc: 0.9801
Epoch 19/20
60000/60000 [=====] -
    7s 114us/step - loss: 0.1083 - acc: 0.9690 -
val_loss: 0.0737 - val_acc: 0.9798
Epoch 20/20
60000/60000 [=====] -
    7s 112us/step - loss: 0.1036 - acc: 0.9700 -
val_loss: 0.0704 - val_acc: 0.9802
```

In [78]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

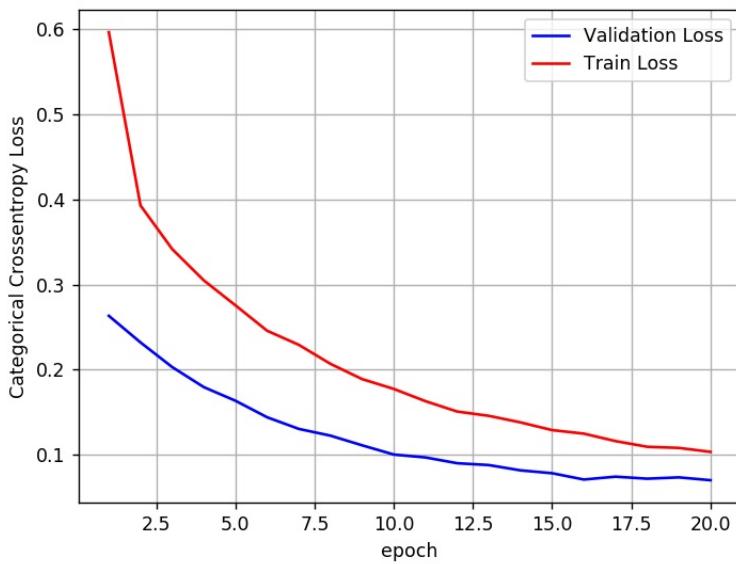
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0703752794698812

Test accuracy: 0.9802



In [79]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

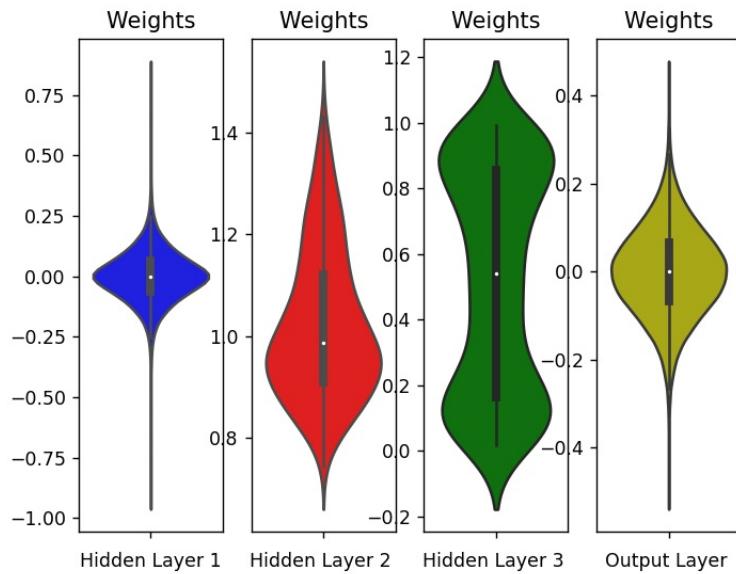
fig = plt.figure()
plt.title("Weight matrices after model trained")

plt.subplot(1, 4, 1)
plt.title("Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 4, 3)
plt.title("Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(1, 4, 4)
plt.title("Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [80]:

```
from keras.optimizers import Adam, RMSprop, SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(256, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)) )
    model.add(Dense(64, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

In [81]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid', 'relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
```

```
model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid,
    n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

In [82]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.975400 using {'activ': 'sigmoid'}
0.975400 (0.001167) with: {'activ': 'sigmoid'}
0.974300 (0.002046) with: {'activ': 'relu'}
```

**architecture 3 -hidden layers 5: input
(28 *28) hidden layer-1 (512) layer-2
(256) hidden layer-3 (128) hidden
layer-4 (64) hidden layer-5 (32) output
softmax 10**

MLP + Sigmoid activation + SGDOptimizer

In [83]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(256, activation='sigmoid'))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(64, activation='sigmoid'))
model_sigmoid.add(Dense(32, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Layer (type)	Output Shape
Param #	
=====	=====
dense_93 (Dense)	(None, 512)

401920

dense_94 (Dense)	(None, 256)
131328	

dense_95 (Dense)	(None, 128)
32896	

dense_96 (Dense)	(None, 64)
8256	

dense_97 (Dense)	(None, 32)
2080	

dense_98 (Dense)	(None, 10)
330	

Total params: 576,810

Trainable params: 576,810

Non-trainable params: 0

In [84]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] -
5s 91us/step - loss: 2.3155 - acc: 0.1124 - val_loss: 2.3012 - val_acc: 0.1135

Epoch 2/20

60000/60000 [=====] -
4s 66us/step - loss: 2.3014 - acc: 0.1124 - val_loss: 2.3012 - val_acc: 0.1135

Epoch 3/20

60000/60000 [=====] -
4s 68us/step - loss: 2.3014 - acc: 0.1124 - val_loss: 2.3009 - val_acc: 0.1135

Epoch 4/20

60000/60000 [=====] -
4s 66us/step - loss: 2.3013 - acc: 0.1124 - val_loss: 2.3009 - val_acc: 0.1135

Epoch 5/20

60000/60000 [=====] -
4s 66us/step - loss: 2.3013 - acc: 0.1124 - val_loss: 2.3011 - val_acc: 0.1135

Epoch 6/20

60000/60000 [=====] -
4s 67us/step - loss: 2.3013 - acc: 0.1124 - val_loss: 2.3010 - val_acc: 0.1135

Epoch 7/20

60000/60000 [=====] -
4s 71us/step - loss: 2.3012 - acc: 0.1124 - val_loss: 2.3008 - val_acc: 0.1135

Epoch 8/20

60000/60000 [=====] -
4s 69us/step - loss: 2.3012 - acc: 0.1124 - val_loss: 2.3011 - val_acc: 0.1135

Epoch 9/20

60000/60000 [=====] -
4s 68us/step - loss: 2.3012 - acc: 0.1124 - val_loss: 2.3011 - val_acc: 0.1135

```
al_loss: 2.3010 - val_acc: 0.1135
Epoch 10/20
60000/60000 [=====] -
4s 68us/step - loss: 2.3012 - acc: 0.1123 - v
al_loss: 2.3009 - val_acc: 0.1135
Epoch 11/20
60000/60000 [=====] -
4s 73us/step - loss: 2.3011 - acc: 0.1124 - v
al_loss: 2.3011 - val_acc: 0.1135
Epoch 12/20
60000/60000 [=====] -
4s 71us/step - loss: 2.3011 - acc: 0.1124 - v
al_loss: 2.3007 - val_acc: 0.1135
Epoch 13/20
60000/60000 [=====] -
4s 71us/step - loss: 2.3010 - acc: 0.1124 - v
al_loss: 2.3010 - val_acc: 0.1135
Epoch 14/20
60000/60000 [=====] -
4s 72us/step - loss: 2.3011 - acc: 0.1124 - v
al_loss: 2.3009 - val_acc: 0.1135
Epoch 15/20
60000/60000 [=====] -
4s 73us/step - loss: 2.3010 - acc: 0.1124 - v
al_loss: 2.3009 - val_acc: 0.1135
Epoch 16/20
60000/60000 [=====] -
4s 71us/step - loss: 2.3010 - acc: 0.1124 - v
al_loss: 2.3008 - val_acc: 0.1135
Epoch 17/20
60000/60000 [=====] -
4s 71us/step - loss: 2.3010 - acc: 0.1124 - v
al_loss: 2.3008 - val_acc: 0.1135
Epoch 18/20
60000/60000 [=====] -
4s 72us/step - loss: 2.3009 - acc: 0.1124 - v
al_loss: 2.3006 - val_acc: 0.1135
```

```
Epoch 19/20
60000/60000 [=====] -
 4s 72us/step - loss: 2.3009 - acc: 0.1124 - val_loss: 2.3006 - val_acc: 0.1135
Epoch 20/20
60000/60000 [=====] -
 4s 71us/step - loss: 2.3009 - acc: 0.1124 - val_loss: 2.3005 - val_acc: 0.1135
```

In [85]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

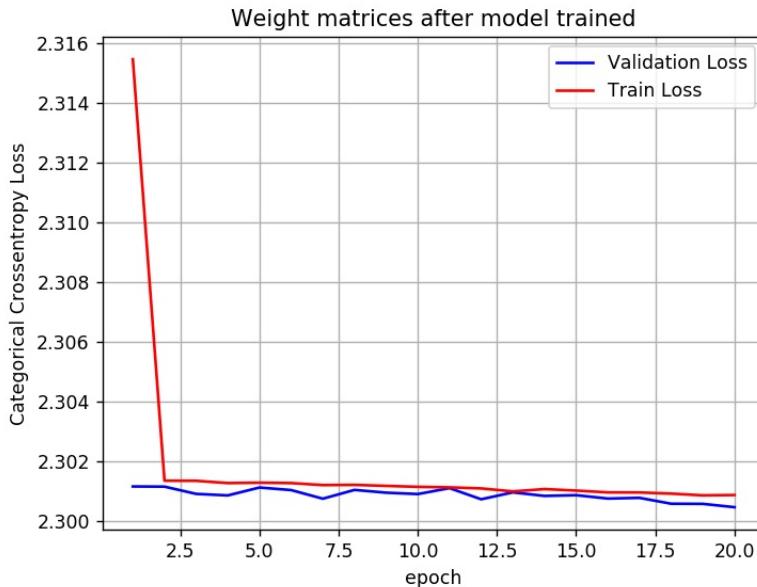
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 2.3004731758117676

Test accuracy: 0.1135



In [93]:

```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

plt.title("Weight matrices after model trained")
fig = plt.figure(figsize=(10,10))
```

```
plt.subplot(2, 3, 1)
plt.title("model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

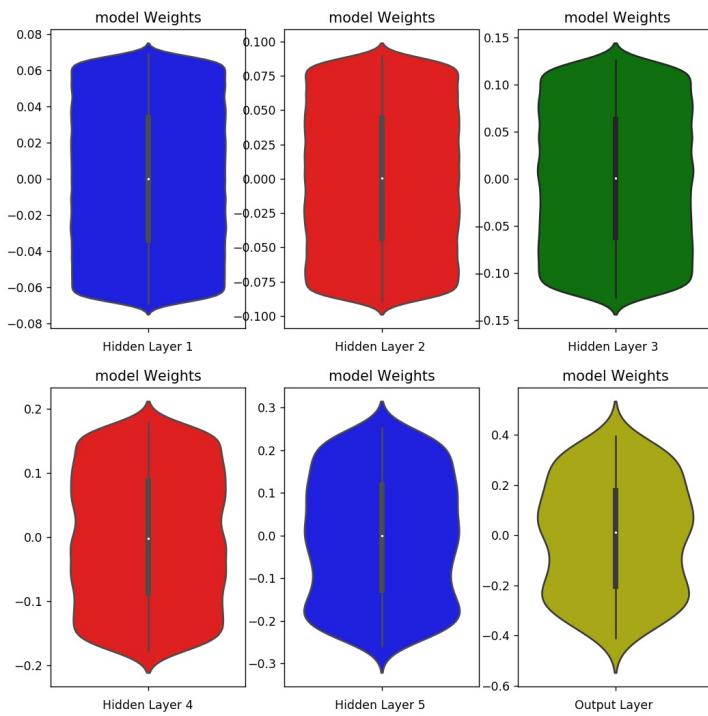
plt.subplot(2, 3, 2)
plt.title("model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 3, 3)
plt.title(" model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("model Weights")
ax = sns.violinplot(y=h4_w,color='r')
plt.xlabel('Hidden Layer 4')

plt.subplot(2, 3, 5)
plt.title("model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2, 3, 6)
plt.title("model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + ADAM

In [94]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(256, activation='sigmoid'))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(64, activation='sigmoid'))
model_sigmoid.add(Dense(32, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape
Param #	
dense_99 (Dense)	(None, 512)
	401920
dense_100 (Dense)	(None, 256)
	131328

dense_101 (Dense)	(None, 128)
32896	

dense_102 (Dense)	(None, 64)
8256	

dense_103 (Dense)	(None, 32)
2080	

dense_104 (Dense)	(None, 10)
330	

Total params: 576,810
Trainable params: 576,810
Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] -
7s 122us/step - loss: 1.2380 - acc: 0.6232 -
val_loss: 0.5369 - val_acc: 0.8743

Epoch 2/20

60000/60000 [=====] -
6s 95us/step - loss: 0.3561 - acc: 0.9166 -
val_loss: 0.2575 - val_acc: 0.9386

Epoch 3/20

60000/60000 [=====] -
6s 92us/step - loss: 0.2012 - acc: 0.9512 -
val_loss: 0.1934 - val_acc: 0.9513

```
Epoch 4/20
60000/60000 [=====] -
  6s 96us/step - loss: 0.1457 - acc: 0.9628 - val_loss: 0.1535 - val_acc: 0.9605
Epoch 5/20
60000/60000 [=====] -
  6s 96us/step - loss: 0.1150 - acc: 0.9700 - val_loss: 0.1309 - val_acc: 0.9659
Epoch 6/20
60000/60000 [=====] -
  6s 92us/step - loss: 0.0915 - acc: 0.9762 - val_loss: 0.1161 - val_acc: 0.9681
Epoch 7/20
60000/60000 [=====] -
  6s 95us/step - loss: 0.0762 - acc: 0.9798 - val_loss: 0.1068 - val_acc: 0.9704
Epoch 8/20
60000/60000 [=====] -
  6s 99us/step - loss: 0.0638 - acc: 0.9833 - val_loss: 0.1004 - val_acc: 0.9716
Epoch 9/20
60000/60000 [=====] -
  6s 96us/step - loss: 0.0531 - acc: 0.9860 - val_loss: 0.0984 - val_acc: 0.9748
Epoch 10/20
60000/60000 [=====] -
  6s 97us/step - loss: 0.0477 - acc: 0.9871 - val_loss: 0.0944 - val_acc: 0.9744
Epoch 11/20
60000/60000 [=====] -
  6s 101us/step - loss: 0.0396 - acc: 0.9896 - val_loss: 0.0867 - val_acc: 0.9776
Epoch 12/20
60000/60000 [=====] -
  6s 97us/step - loss: 0.0348 - acc: 0.9905 - val_loss: 0.0888 - val_acc: 0.9755
Epoch 13/20
```

```
60000/60000 [=====] -  
 6s 98us/step - loss: 0.0285 - acc: 0.9925 - v  
al_loss: 0.0896 - val_acc: 0.9771  
Epoch 14/20  
60000/60000 [=====] -  
 6s 96us/step - loss: 0.0250 - acc: 0.9931 - v  
al_loss: 0.0965 - val_acc: 0.9744  
Epoch 15/20  
60000/60000 [=====] -  
 6s 94us/step - loss: 0.0238 - acc: 0.9933 - v  
al_loss: 0.0902 - val_acc: 0.9771  
Epoch 16/20  
60000/60000 [=====] -  
 6s 96us/step - loss: 0.0195 - acc: 0.9947 - v  
al_loss: 0.0886 - val_acc: 0.9771  
Epoch 17/20  
60000/60000 [=====] -  
 6s 93us/step - loss: 0.0183 - acc: 0.9952 - v  
al_loss: 0.0999 - val_acc: 0.9750  
Epoch 18/20  
60000/60000 [=====] -  
 6s 94us/step - loss: 0.0152 - acc: 0.9958 - v  
al_loss: 0.0910 - val_acc: 0.9790  
Epoch 19/20  
60000/60000 [=====] -  
 6s 100us/step - loss: 0.0161 - acc: 0.9957 -  
val_loss: 0.0896 - val_acc: 0.9798  
Epoch 20/20  
60000/60000 [=====] -  
 6s 98us/step - loss: 0.0116 - acc: 0.9967 - v  
al_loss: 0.0967 - val_acc: 0.9781
```

In [95]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])
```

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

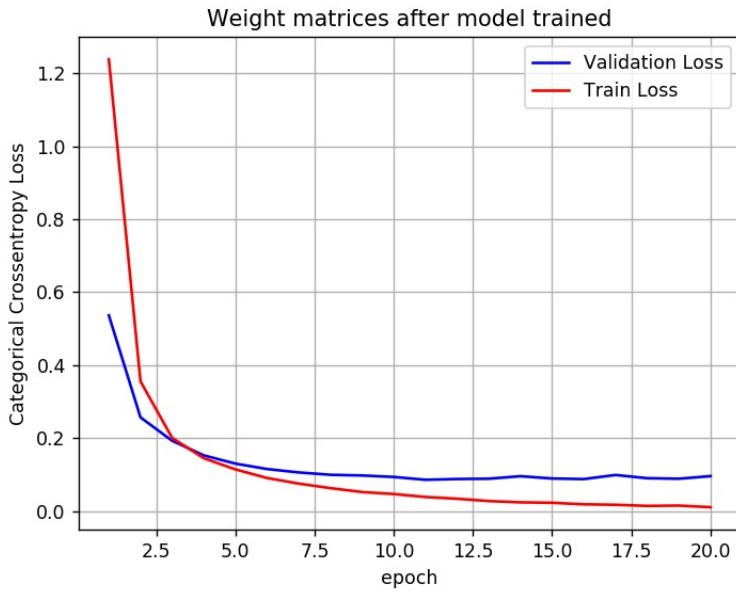
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09669930382110178

Test accuracy: 0.9781



In [96]:

```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

plt.title("Weight matrices after model trained")
fig = plt.figure(figsize=(10,10))

plt.subplot(2, 3, 1)
plt.title("model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("model Weights")
```

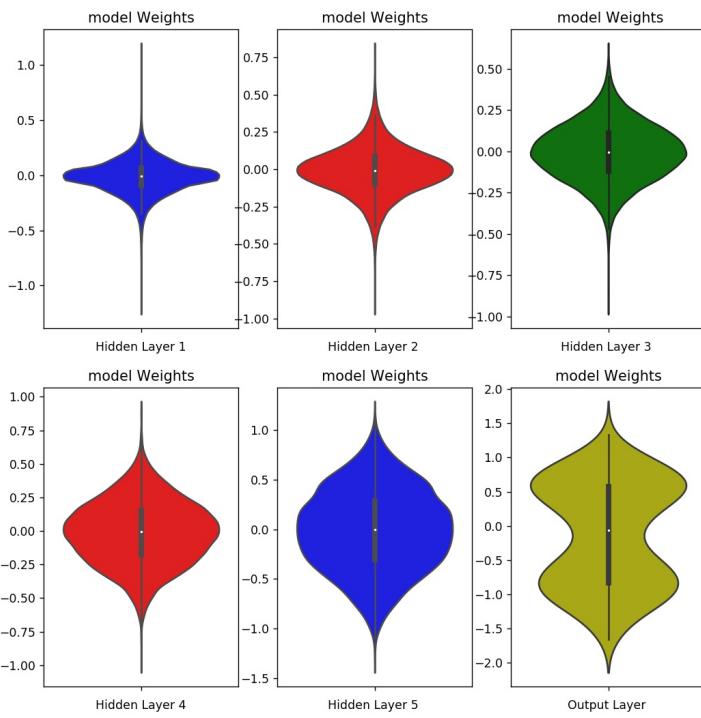
```
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(2, 3, 3)
plt.title(" model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(2, 3, 4)
plt.title("model Weights")
ax = sns.violinplot(y=h4_w,color='r')
plt.xlabel('Hidden Layer 4')

plt.subplot(2, 3, 5)
plt.title("model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5')

plt.subplot(2, 3, 6)
plt.title("model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU +SGD

In [97]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(\theta, \sigma) = N(\theta, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.088 \Rightarrow N(\theta, \sigma) = N(\theta, 0.088)$ 
# h3 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(\theta, \sigma) = N(\theta, 0.125)$ 
# h4 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.176 \Rightarrow N(\theta, \sigma) = N(\theta, 0.176)$ 
# h5 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.25 \Rightarrow N(\theta, \sigma) = N(\theta, 0.25)$ 

# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.218 \Rightarrow N(\theta, \sigma) = N(\theta, 0.218)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)) )
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)) )
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.25, seed=None)) )

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape
Param #	
=====	=====
dense_105 (Dense)	(None, 512)
401920	
=====	=====
dense_106 (Dense)	(None, 256)
131328	
=====	=====
dense_107 (Dense)	(None, 128)
32896	
=====	=====
dense_108 (Dense)	(None, 64)
8256	
=====	=====
dense_109 (Dense)	(None, 32)
2080	
=====	=====
dense_110 (Dense)	(None, 10)
330	
=====	=====
Total params:	576,810
Trainable params:	576,810
Non-trainable params:	0

In [98]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

```
60000/60000 [=====] -  
 6s 94us/step - loss: 0.5969 - acc: 0.8209 - val_loss: 0.2808 - val_acc: 0.9167
```

Epoch 2/20

```
60000/60000 [=====] -  
 4s 67us/step - loss: 0.2519 - acc: 0.9246 - val_loss: 0.2160 - val_acc: 0.9363
```

Epoch 3/20

```
60000/60000 [=====] -  
 4s 67us/step - loss: 0.1956 - acc: 0.9420 - val_loss: 0.1938 - val_acc: 0.9420
```

Epoch 4/20

```
60000/60000 [=====] -  
 4s 68us/step - loss: 0.1647 - acc: 0.9514 - val_loss: 0.1948 - val_acc: 0.9415
```

Epoch 5/20

```
60000/60000 [=====] -  
 4s 68us/step - loss: 0.1431 - acc: 0.9576 - val_loss: 0.1582 - val_acc: 0.9490
```

Epoch 6/20

```
60000/60000 [=====] -  
 4s 69us/step - loss: 0.1254 - acc: 0.9627 - val_loss: 0.1460 - val_acc: 0.9542
```

Epoch 7/20

```
60000/60000 [=====] -  
 4s 69us/step - loss: 0.1126 - acc: 0.9663 - val_loss: 0.1500 - val_acc: 0.9543
```

Epoch 8/20
60000/60000 [=====] -
4s 72us/step - loss: 0.1018 - acc: 0.9697 - v
al_loss: 0.1372 - val_acc: 0.9592

Epoch 9/20
60000/60000 [=====] -
4s 72us/step - loss: 0.0905 - acc: 0.9738 - v
al_loss: 0.1278 - val_acc: 0.9618

Epoch 10/20
60000/60000 [=====] -
4s 72us/step - loss: 0.0829 - acc: 0.9758 - v
al_loss: 0.1254 - val_acc: 0.9629

Epoch 11/20
60000/60000 [=====] -
4s 69us/step - loss: 0.0753 - acc: 0.9781 - v
al_loss: 0.1187 - val_acc: 0.9640

Epoch 12/20
60000/60000 [=====] -
4s 73us/step - loss: 0.0690 - acc: 0.9803 - v
al_loss: 0.1258 - val_acc: 0.9618

Epoch 13/20
60000/60000 [=====] -
4s 71us/step - loss: 0.0631 - acc: 0.9821 - v
al_loss: 0.1145 - val_acc: 0.9659

Epoch 14/20
60000/60000 [=====] -
4s 72us/step - loss: 0.0579 - acc: 0.9840 - v
al_loss: 0.1198 - val_acc: 0.9647

Epoch 15/20
60000/60000 [=====] -
4s 72us/step - loss: 0.0526 - acc: 0.9856 - v
al_loss: 0.1170 - val_acc: 0.9646

Epoch 16/20
60000/60000 [=====] -
4s 75us/step - loss: 0.0484 - acc: 0.9869 - v
al_loss: 0.1135 - val_acc: 0.9660

Epoch 17/20

```
60000/60000 [=====] -  
 4s 72us/step - loss: 0.0444 - acc: 0.9880 - v  
al_loss: 0.1133 - val_acc: 0.9671  
Epoch 18/20  
60000/60000 [=====] -  
 4s 71us/step - loss: 0.0411 - acc: 0.9890 - v  
al_loss: 0.1095 - val_acc: 0.9684  
Epoch 19/20  
60000/60000 [=====] -  
 4s 73us/step - loss: 0.0371 - acc: 0.9906 - v  
al_loss: 0.1139 - val_acc: 0.9671  
Epoch 20/20  
60000/60000 [=====] -  
 4s 73us/step - loss: 0.0344 - acc: 0.9914 - v  
al_loss: 0.1073 - val_acc: 0.9679
```

In [99]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig,ax = plt.subplots(1,1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')  
  
# list of epoch numbers  
x = list(range(1,nb_epoch+1))  
  
# print(history.history.keys())  
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])  
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))  
  
# we will get val_loss and val_acc only when you pass the parameter validation_data
```

```

# val_loss : validation loss
# val_acc : validation accuracy

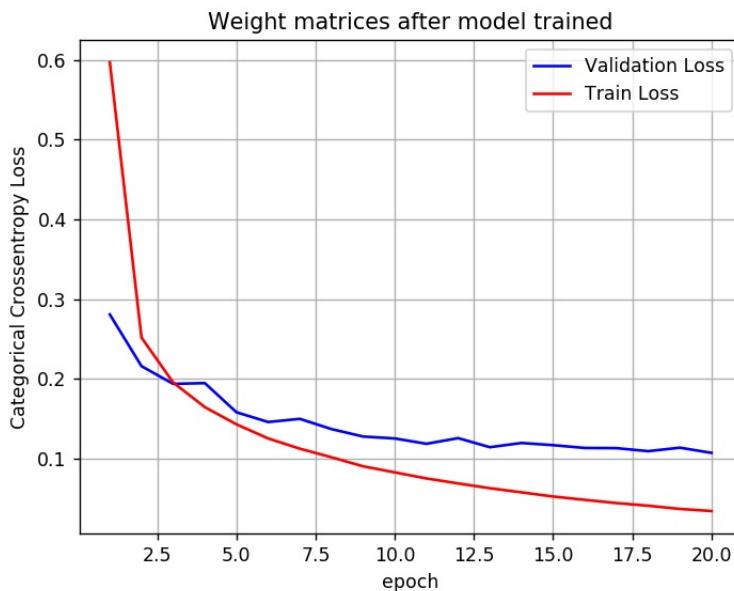
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.1073460508538934

Test accuracy: 0.9679



In [100]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)

```

```
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

plt.title("Weight matrices after model trained")
fig = plt.figure(figsize=(10,10))

plt.subplot(2, 3, 1)
plt.title("model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

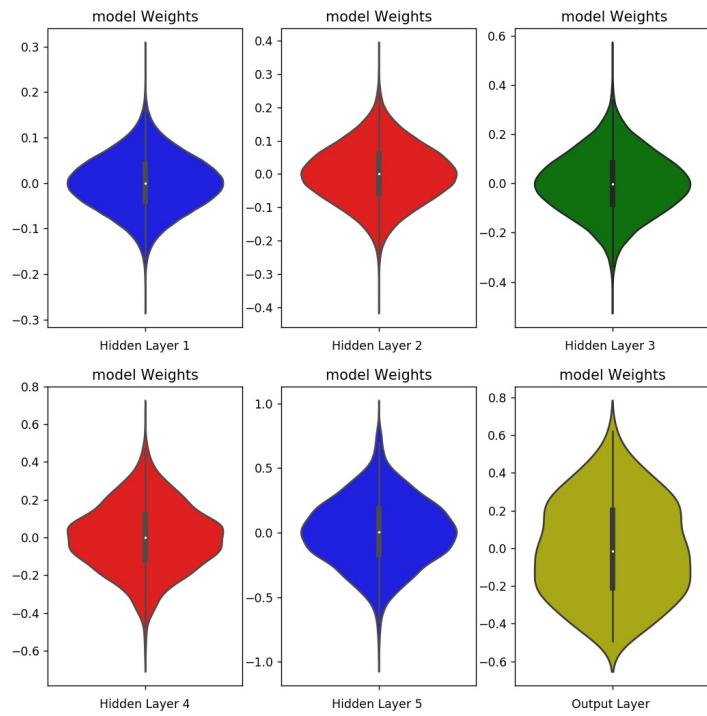
plt.subplot(2, 3, 3)
plt.title(" model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("model Weights")
ax = sns.violinplot(y=h4_w,color='r')
plt.xlabel('Hidden Layer 4')

plt.subplot(2, 3, 5)
plt.title("model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2, 3, 6)
plt.title("model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
```

```
plt.show()
```



MLP + ReLU + ADAM

In [101]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None) ) )
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)) )
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.25, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))


print(model_relu.summary())


model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape
Param #	
=====	=====
dense_111 (Dense)	(None, 512)

401920

dense_112 (Dense)	(None, 256)
131328	

dense_113 (Dense)	(None, 128)
32896	

dense_114 (Dense)	(None, 64)
8256	

dense_115 (Dense)	(None, 32)
2080	

dense_116 (Dense)	(None, 10)
330	

Total params: 576,810

Trainable params: 576,810

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] -

7s 123us/step - loss: 0.2748 - acc: 0.9183 -

val_loss: 0.1250 - val_acc: 0.9633

Epoch 2/20

60000/60000 [=====] -

```
6s 96us/step - loss: 0.1003 - acc: 0.9698 - v  
al_loss: 0.0967 - val_acc: 0.9696  
Epoch 3/20  
60000/60000 [=====] -  
6s 97us/step - loss: 0.0657 - acc: 0.9797 - v  
al_loss: 0.0949 - val_acc: 0.9714  
Epoch 4/20  
60000/60000 [=====] -  
6s 94us/step - loss: 0.0468 - acc: 0.9853 - v  
al_loss: 0.0793 - val_acc: 0.9767  
Epoch 5/20  
60000/60000 [=====] -  
6s 97us/step - loss: 0.0392 - acc: 0.9875 - v  
al_loss: 0.0859 - val_acc: 0.9756  
Epoch 6/20  
60000/60000 [=====] -  
6s 100us/step - loss: 0.0338 - acc: 0.9889 -  
val_loss: 0.0885 - val_acc: 0.9764  
Epoch 7/20  
60000/60000 [=====] -  
6s 101us/step - loss: 0.0271 - acc: 0.9917 -  
val_loss: 0.0940 - val_acc: 0.9763  
Epoch 8/20  
60000/60000 [=====] -  
6s 103us/step - loss: 0.0262 - acc: 0.9915 -  
val_loss: 0.0992 - val_acc: 0.9767  
Epoch 9/20  
60000/60000 [=====] -  
6s 99us/step - loss: 0.0252 - acc: 0.9924 - v  
al_loss: 0.1119 - val_acc: 0.9709  
Epoch 10/20  
60000/60000 [=====] -  
6s 98us/step - loss: 0.0199 - acc: 0.9935 - v  
al_loss: 0.0861 - val_acc: 0.9802  
Epoch 11/20  
60000/60000 [=====] -  
6s 103us/step - loss: 0.0166 - acc: 0.9946 -
```

```
val_loss: 0.0939 - val_acc: 0.9772
Epoch 12/20
60000/60000 [=====] -
  6s 105us/step - loss: 0.0195 - acc: 0.9941 -
val_loss: 0.0975 - val_acc: 0.9764
Epoch 13/20
60000/60000 [=====] -
  6s 99us/step - loss: 0.0182 - acc: 0.9941 - v
al_loss: 0.0843 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] -
  6s 96us/step - loss: 0.0162 - acc: 0.9950 - v
al_loss: 0.1064 - val_acc: 0.9778
Epoch 15/20
60000/60000 [=====] -
  6s 96us/step - loss: 0.0142 - acc: 0.9957 - v
al_loss: 0.0859 - val_acc: 0.9812
Epoch 16/20
60000/60000 [=====] -
  6s 99us/step - loss: 0.0148 - acc: 0.9956 - v
al_loss: 0.1164 - val_acc: 0.9743
Epoch 17/20
60000/60000 [=====] -
  6s 97us/step - loss: 0.0145 - acc: 0.9956 - v
al_loss: 0.1061 - val_acc: 0.9784
Epoch 18/20
60000/60000 [=====] -
  6s 97us/step - loss: 0.0098 - acc: 0.9968 - v
al_loss: 0.1287 - val_acc: 0.9743
Epoch 19/20
60000/60000 [=====] -
  6s 99us/step - loss: 0.0133 - acc: 0.9960 - v
al_loss: 0.1151 - val_acc: 0.9762
Epoch 20/20
60000/60000 [=====] -
  6s 96us/step - loss: 0.0133 - acc: 0.9961 - v
al_loss: 0.1028 - val_acc: 0.9778
```

In [102]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10280663507273835

Test accuracy: 0.9778



In [103]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

plt.title("Weight matrices after model trained")
fig = plt.figure(figsize=(10,10))

plt.subplot(2, 3, 1)
plt.title("model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("model Weights")
```

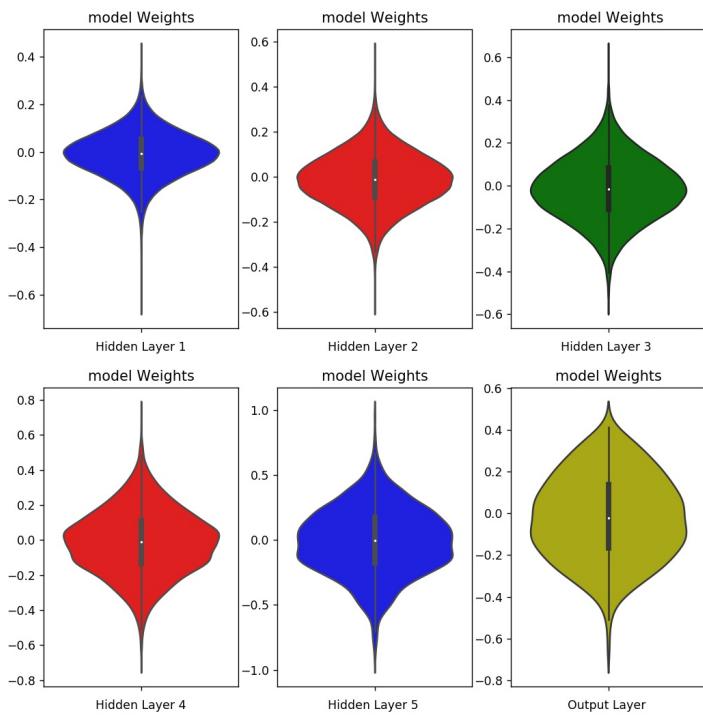
```
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(2, 3, 3)
plt.title(" model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(2, 3, 4)
plt.title("model Weights")
ax = sns.violinplot(y=h4_w,color='r')
plt.xlabel('Hidden Layer 4')

plt.subplot(2, 3, 5)
plt.title("model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5')

plt.subplot(2, 3, 6)
plt.title("model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [104]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.051 \Rightarrow N(\theta, \sigma) = N(0, 0.051)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.072 \Rightarrow N(\theta, \sigma) = N(0, 0.072)$ 
# h4 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.102 \Rightarrow N(\theta, \sigma) = N(0, 0.102)$ 
# h5 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.144 \Rightarrow N(\theta, \sigma) = N(0, 0.144)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None)))
model_batch.add(BatchNormalization())
```

```

model_batch.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.102, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(32, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.144, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape
Param #	
=====	=====
dense_117 (Dense)	(None, 512)
401920	
=====	=====
batch_normalization_11 (BatchNormalization)	(None, 512)
2048	
=====	=====
dense_118 (Dense)	(None, 256)
131328	
=====	=====
batch_normalization_12 (BatchNormalization)	(None, 256)
1024	
=====	=====
dense_119 (Dense)	(None, 128)
32896	

```
batch_normalization_13 (Batch Normalization) (None, 128)
512

dense_120 (Dense) (None, 64)
8256

batch_normalization_14 (Batch Normalization) (None, 64)
256

dense_121 (Dense) (None, 32)
2080

batch_normalization_15 (Batch Normalization) (None, 32)
128

dense_122 (Dense) (None, 10)
330
=====
=====
```

Total params: 580,778
Trainable params: 578,794
Non-trainable params: 1,984

In [105]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/20
```

```
60000/60000 [=====] -  
18s 295us/step - loss: 0.2675 - acc: 0.9238 -  
val_loss: 0.1430 - val_acc: 0.9581
```

```
Epoch 2/20
```

```
60000/60000 [=====] -  
12s 208us/step - loss: 0.1196 - acc: 0.9640 -  
val_loss: 0.1144 - val_acc: 0.9642
```

```
Epoch 3/20
```

```
60000/60000 [=====] -  
13s 213us/step - loss: 0.0864 - acc: 0.9735 -  
val_loss: 0.1112 - val_acc: 0.9659
```

```
Epoch 4/20
```

```
60000/60000 [=====] -  
13s 220us/step - loss: 0.0644 - acc: 0.9797 -  
val_loss: 0.0830 - val_acc: 0.9741
```

```
Epoch 5/20
```

```
60000/60000 [=====] -  
14s 226us/step - loss: 0.0527 - acc: 0.9833 -  
val_loss: 0.0840 - val_acc: 0.9732
```

```
Epoch 6/20
```

```
60000/60000 [=====] -  
13s 220us/step - loss: 0.0457 - acc: 0.9856 -  
val_loss: 0.0887 - val_acc: 0.9721
```

```
Epoch 7/20
```

```
60000/60000 [=====] -  
14s 227us/step - loss: 0.0365 - acc: 0.9883 -  
val_loss: 0.0791 - val_acc: 0.9766
```

```
Epoch 8/20
```

```
60000/60000 [=====] -  
13s 212us/step - loss: 0.0322 - acc: 0.9893 -
```

```
val_loss: 0.0875 - val_acc: 0.9757
Epoch 9/20
60000/60000 [=====] -
12s 201us/step - loss: 0.0287 - acc: 0.9905 -
val_loss: 0.0930 - val_acc: 0.9749
Epoch 10/20
60000/60000 [=====] -
11s 189us/step - loss: 0.0255 - acc: 0.9915 -
val_loss: 0.0926 - val_acc: 0.9758
Epoch 11/20
60000/60000 [=====] -
13s 218us/step - loss: 0.0277 - acc: 0.9911 -
val_loss: 0.0770 - val_acc: 0.9809
Epoch 12/20
60000/60000 [=====] -
12s 205us/step - loss: 0.0197 - acc: 0.9932 -
val_loss: 0.0855 - val_acc: 0.9788
Epoch 13/20
60000/60000 [=====] -
13s 212us/step - loss: 0.0204 - acc: 0.9932 -
val_loss: 0.0916 - val_acc: 0.9772
Epoch 14/20
60000/60000 [=====] -
13s 213us/step - loss: 0.0178 - acc: 0.9941 -
val_loss: 0.0915 - val_acc: 0.9756
Epoch 15/20
60000/60000 [=====] -
13s 217us/step - loss: 0.0212 - acc: 0.9926 -
val_loss: 0.0852 - val_acc: 0.9766
Epoch 16/20
60000/60000 [=====] -
13s 212us/step - loss: 0.0183 - acc: 0.9940 -
val_loss: 0.1004 - val_acc: 0.9769
Epoch 17/20
60000/60000 [=====] -
12s 200us/step - loss: 0.0154 - acc: 0.9949 -
val_loss: 0.0848 - val_acc: 0.9779
```

```
Epoch 18/20
60000/60000 [=====] -
 13s 217us/step - loss: 0.0151 - acc: 0.9950 -
 val_loss: 0.0910 - val_acc: 0.9794
Epoch 19/20
60000/60000 [=====] -
 13s 213us/step - loss: 0.0160 - acc: 0.9946 -
 val_loss: 0.0862 - val_acc: 0.9791
Epoch 20/20
60000/60000 [=====] -
 11s 183us/step - loss: 0.0134 - acc: 0.9953 -
 val_loss: 0.0891 - val_acc: 0.9780
```

In [106]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
```

```

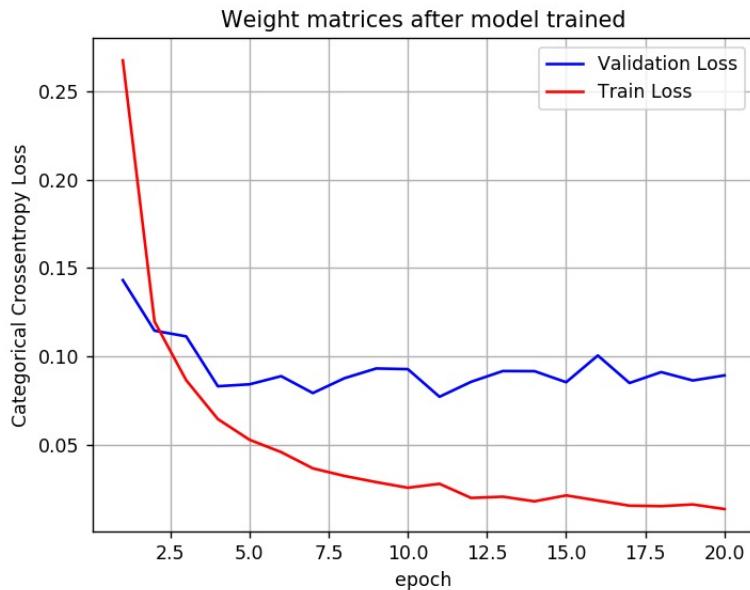
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08909254919743981

Test accuracy: 0.978



In [107]:

```

w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

```

```
plt.title("Weight matrices after model trained")
fig = plt.figure(figsize=(10,10))

plt.subplot(2, 3, 1)
plt.title("model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

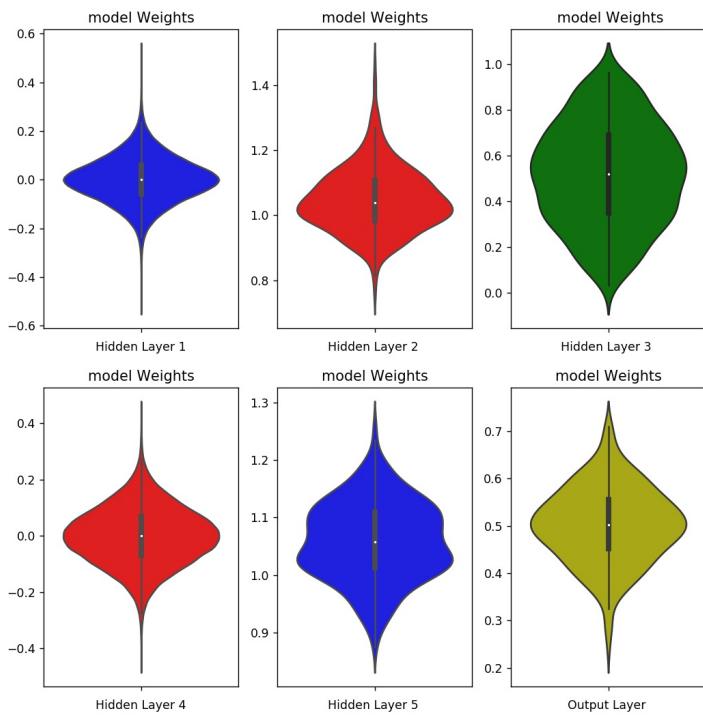
plt.subplot(2, 3, 2)
plt.title("model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 3, 3)
plt.title(" model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("model Weights")
ax = sns.violinplot(y=h4_w,color='r')
plt.xlabel('Hidden Layer 4')

plt.subplot(2, 3, 5)
plt.title("model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2, 3, 6)
plt.title("model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5. MLP + Dropout + AdamOptimizer

In [108]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None) ))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None) ))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.102, seed=None) ))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.144, seed=None) ))
```

```
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape
Param #	
dense_123 (Dense)	(None, 512)
	401920
batch_normalization_16 (Batch Normalization)	(None, 512)
	2048
dropout_6 (Dropout)	(None, 512)
	0
dense_124 (Dense)	(None, 256)
	131328
batch_normalization_17 (Batch Normalization)	(None, 256)
	1024
dropout_7 (Dropout)	(None, 256)
	0

dense_125 (Dense) (None, 128)
32896

batch_normalization_18 (BatchNormalization) (None, 128)
512

dropout_8 (Dropout) (None, 128)
0

dense_126 (Dense) (None, 64)
8256

batch_normalization_19 (BatchNormalization) (None, 64)
256

dropout_9 (Dropout) (None, 64)
0

dense_127 (Dense) (None, 32)
2080

batch_normalization_20 (BatchNormalization) (None, 32)
128

dropout_10 (Dropout) (None, 32)
0

```
dense_128 (Dense)           (None, 10)
    330
=====
=====
Total params: 580,778
Trainable params: 578,794
Non-trainable params: 1,984
```

In [109]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy',
                    metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
                          epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

```
60000/60000 [=====] -
 12s 193us/step - loss: 1.0613 - acc: 0.6596 -
 val_loss: 0.3441 - val_acc: 0.9018
```

Epoch 2/20

```
60000/60000 [=====] -
 9s 146us/step - loss: 0.6044 - acc: 0.8298 -
 val_loss: 0.2848 - val_acc: 0.9183
```

Epoch 3/20

```
60000/60000 [=====] -
 9s 147us/step - loss: 0.5054 - acc: 0.8617 -
 val_loss: 0.2638 - val_acc: 0.9269
```

Epoch 4/20

```
60000/60000 [=====] -
 9s 151us/step - loss: 0.4489 - acc: 0.8819 -
 val_loss: 0.2435 - val_acc: 0.9322
```

Epoch 5/20

```
60000/60000 [=====] -  
 9s 155us/step - loss: 0.3990 - acc: 0.8964 -  
val_loss: 0.2117 - val_acc: 0.9418  
Epoch 6/20  
60000/60000 [=====] -  
 9s 149us/step - loss: 0.3553 - acc: 0.9095 -  
val_loss: 0.1913 - val_acc: 0.9503  
Epoch 7/20  
60000/60000 [=====] -  
 9s 147us/step - loss: 0.3234 - acc: 0.9182 -  
val_loss: 0.1709 - val_acc: 0.9562  
Epoch 8/20  
60000/60000 [=====] -  
 9s 147us/step - loss: 0.2942 - acc: 0.9258 -  
val_loss: 0.1582 - val_acc: 0.9599  
Epoch 9/20  
60000/60000 [=====] -  
 9s 148us/step - loss: 0.2726 - acc: 0.9328 -  
val_loss: 0.1441 - val_acc: 0.9645  
Epoch 10/20  
60000/60000 [=====] -  
 9s 147us/step - loss: 0.2586 - acc: 0.9378 -  
val_loss: 0.1411 - val_acc: 0.9665  
Epoch 11/20  
60000/60000 [=====] -  
 9s 148us/step - loss: 0.2385 - acc: 0.9432 -  
val_loss: 0.1312 - val_acc: 0.9689  
Epoch 12/20  
60000/60000 [=====] -  
 9s 148us/step - loss: 0.2347 - acc: 0.9444 -  
val_loss: 0.1324 - val_acc: 0.9695  
Epoch 13/20  
60000/60000 [=====] -  
 9s 146us/step - loss: 0.2170 - acc: 0.9475 -  
val_loss: 0.1232 - val_acc: 0.9721  
Epoch 14/20  
60000/60000 [=====] -
```

```
9s 150us/step - loss: 0.2130 - acc: 0.9496 -
val_loss: 0.1218 - val_acc: 0.9735
Epoch 15/20
60000/60000 [=====] -
9s 153us/step - loss: 0.2031 - acc: 0.9531 -
val_loss: 0.1176 - val_acc: 0.9741
Epoch 16/20
60000/60000 [=====] -
9s 153us/step - loss: 0.1918 - acc: 0.9547 -
val_loss: 0.1124 - val_acc: 0.9744
Epoch 17/20
60000/60000 [=====] -
9s 147us/step - loss: 0.1850 - acc: 0.9564 -
val_loss: 0.1091 - val_acc: 0.9766
Epoch 18/20
60000/60000 [=====] -
9s 147us/step - loss: 0.1793 - acc: 0.9587 -
val_loss: 0.1082 - val_acc: 0.9772
Epoch 19/20
60000/60000 [=====] -
9s 150us/step - loss: 0.1771 - acc: 0.9585 -
val_loss: 0.1040 - val_acc: 0.9776
Epoch 20/20
60000/60000 [=====] -
9s 152us/step - loss: 0.1662 - acc: 0.9616 -
val_loss: 0.0961 - val_acc: 0.9804
```

In [110]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

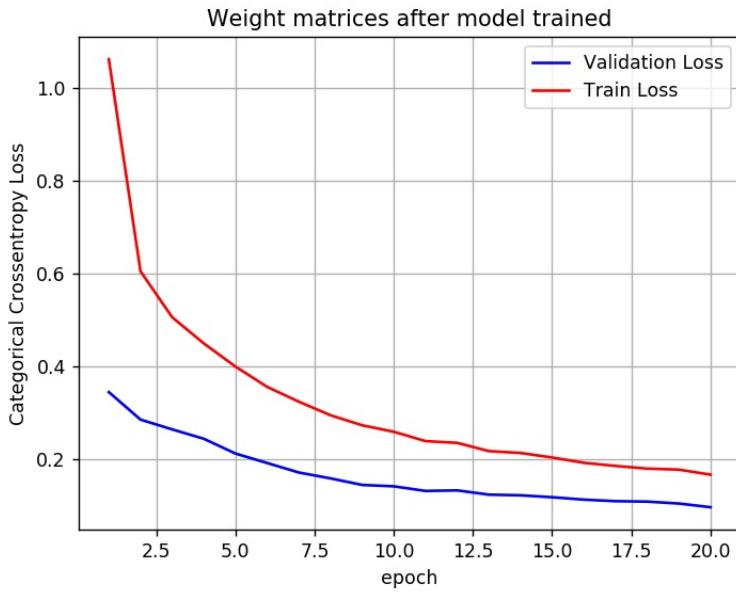
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09613309992719442

Test accuracy: 0.9804



In [111]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

plt.title("Weight matrices after model trained")
fig = plt.figure(figsize=(10,10))

plt.subplot(2, 3, 1)
plt.title("model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 3, 2)
plt.title("model Weights")
```

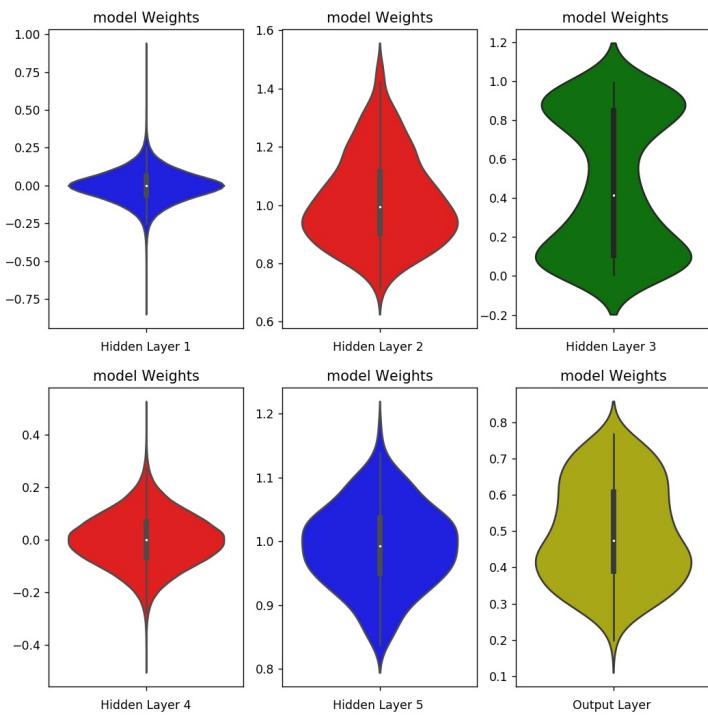
```
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(2, 3, 3)
plt.title(" model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(2, 3, 4)
plt.title("model Weights")
ax = sns.violinplot(y=h4_w,color='r')
plt.xlabel('Hidden Layer 4')

plt.subplot(2, 3, 5)
plt.title("model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5')

plt.subplot(2, 3, 6)
plt.title("model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [112]:

```
from keras.optimizers import Adam, RMSprop, SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(256, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)) )
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(64, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)) )
    model.add(Dense(32, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.25, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

In [113]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/
activ = ['sigmoid', 'relu']
```

```
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs
=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid,
    n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

In [114]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_re
sult.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.974850 using {'activ': 'relu'}
0.971717 (0.000613) with: {'activ': 'sigmoid'}
0.974850 (0.002069) with: {'activ': 'relu'}
```

conclusion

In [116]:

```
# http://zetcode.com/python/prettytable/  
  
from prettytable import PrettyTable  
  
#If you get a ModuleNotFoundError error , install prettytable  
#using: pip3 install prettytable  
  
x = PrettyTable()  
x.field_names = ["MLP architecture ", "Activation function",  
"Optimizer", "Regularization", "Train accuracy", "Test Accuracy"  
]  
  
x.add_row(["Architecture 1", "Sigmoid", "sgd", "None", 0.8705  
, 0.8783])  
x.add_row(["Architecture 1", "Sigmoid", "Adam", "None", 0.999  
4, 0.9828])  
x.add_row(["Architecture 1", "ReLU", "sgd", "None", 0.9699, 0  
.9630])  
x.add_row(["Architecture 1", "ReLU", "Adam", "None", 0.9971,  
0.9810])  
x.add_row(["Architecture 1", "Sigmoid", "Adam", "Batch Normal  
ization", 0.9951, 0.9741])  
x.add_row(["Architecture 1", "Sigmoid", "Adam", "Drop out", 0  
.9470, 0.9679])  
  
x.add_row(["Architecture 2", "Sigmoid", "sgd", "None", 0.6421  
, 0.6652])  
x.add_row(["Architecture 2", "Sigmoid", "Adam", "None", 0.998  
3, 0.9812])
```

```

x.add_row(["Architecture 2", "ReLU", "sgd", "None", 0.9838, 0
.9685])
x.add_row(["Architecture 2", "ReLU", "Adam", "None", 0.9957,
0.9807])
x.add_row(["Architecture 2", "Sigmoid", "Adam", "Batch Normal
ization", 0.9957, 0.9790])
x.add_row(["Architecture 2", "Sigmoid", "Adam", "Drop out", 0
.9700, 0.9802])

x.add_row(["Architecture 3", "Sigmoid", "sgd", "None", 0.1124
, 0.1135])
x.add_row(["Architecture 3", "Sigmoid", "Adam", "None", 0.996
7, 0.9781])
x.add_row(["Architecture 3", "ReLU", "sgd", "None", 0.9914, 0
.9679])
x.add_row(["Architecture 3", "ReLU", "Adam", "None", 0.9961,
0.9778])
x.add_row(["Architecture 3", "Sigmoid", "Adam", "Batch Normal
ization", 0.9953, 0.9780])
x.add_row(["Architecture 3", "Sigmoid", "Adam", "Drop out", 0
.9616, 0.9804])

print(x)

```

MLP architecture	Activation function	Optimizer	Regularization	Train accuracy
Test Accuracy				
Architecture 1	Sigmoid	sgd	None	0.8705
0.8783				

	Architecture 1	Sigmoid	
Adam	None	0.9994	
	0.9828		
	Architecture 1	ReLU	
sgd	None	0.9699	
	0.963		
	Architecture 1	ReLU	
Adam	None	0.9971	
	0.981		
	Architecture 1	Sigmoid	
Adam	Batch Normalization	0.9951	
	0.9741		
	Architecture 1	Sigmoid	
Adam	Drop out	0.947	
	0.9679		
	Architecture 2	Sigmoid	
sgd	None	0.6421	
	0.6652		
	Architecture 2	Sigmoid	
Adam	None	0.9983	
	0.9812		
	Architecture 2	ReLU	
sgd	None	0.9838	
	0.9685		
	Architecture 2	ReLU	
Adam	None	0.9957	
	0.9807		
	Architecture 2	Sigmoid	
Adam	Batch Normalization	0.9957	
	0.979		
	Architecture 2	Sigmoid	
Adam	Drop out	0.97	
	0.9802		
	Architecture 2	Sigmoid	
sgd	None	0.1124	
	0.1135		
	Architecture 3	Sigmoid	

Adam		None		0.9967
	0.9781			
	Architecture 3		ReLU	
sgd		None		0.9914
	0.9679			
	Architecture 3		ReLU	
Adam		None		0.9961
	0.9778			
	Architecture 3		Sigmoid	
Adam		Batch Normalization		0.9953
	0.978			
	Architecture 3		Sigmoid	
Adam		Drop out		0.9616
	0.9804			
-----+-----+-----				
-----+-----+-----				
-----+-----+-----				

Architecture 1 consist of 2 hidden layers with 512 and 128 nuerons respectively. Architecture 2 consist of 3 hidden layers consist of 512, 256 and 64 nuerons respectively. Architecture 3 consist of 5 hidden layers with 512,256,128,64 and 32 nuerons respectively.

conclusion

- 1.We can observe that as hidden layers increases ,the sgd algorithm with sigmoid activation is performing very bad whereas the sgd optimizer with sigmoid do not change significantly.
- 2.With increase in hidden layers the execcution time is increasing
- 3.we do not see significant changes in the train and test accuracy for activation function ReLU and for both optimizers with increase in the hidden layers
- 4.With increase in the no of hidden layers we do not see significant differences in the train and test accuracies.

5. It would be better to take less no of hidden layers while tuning the DNN network.